

# Data Visualization in R

**ggplot2 and the grammar of graphics**

2019-08-29

# ggplot2:

Build a data  
MASTERPIECE



HORST '18

Art by Allison Horst

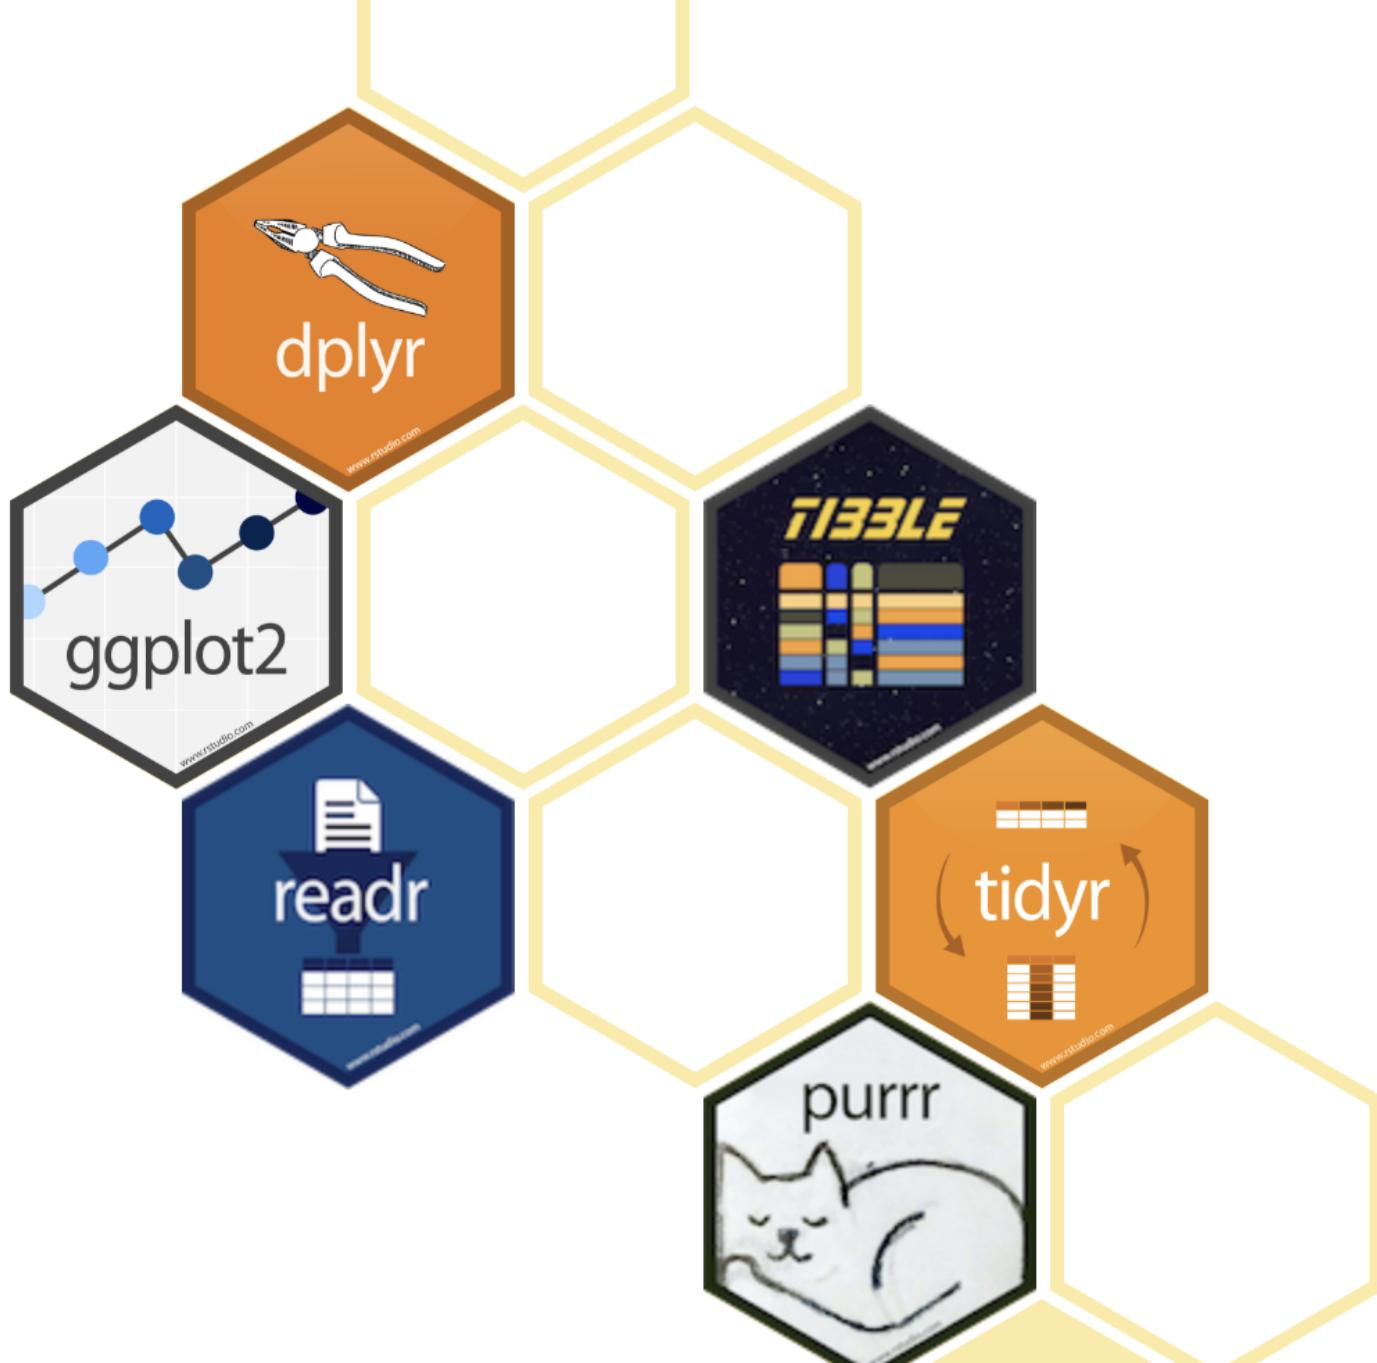
# Data Visualization with R

**ggplot2** works well with the **tidyverse**  
and is friendly and powerful

# Data Visualization with R

**ggplot2 works well with the tidyverse  
and is friendly and powerful**

**Better plots are better communication**



# ggplot2: Elegant Data Visualizations in R

a Layered Grammar of Graphics



# ggplot2: Elegant Data Visualizations in R



a Layered Grammar of Graphics

Data is mapped to aesthetics; Statistics and plot are linked

# ggplot2: Elegant Data Visualizations in R



a Layered Grammar of Graphics

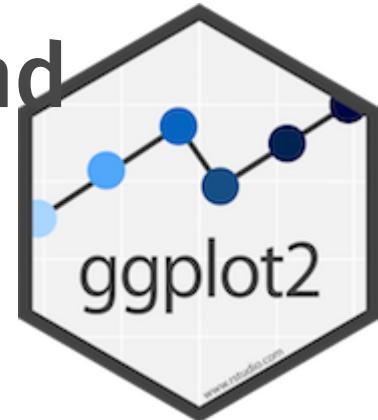
Data is mapped to aesthetics; Statistics and plot are linked

Sensible defaults; Infinitely extensible

# Publication quality and beyond

<https://nyti.ms/2jUp36n>

<http://bit.ly/2KSGZLu>



```
# print prettily  
as_tibble(mtcars)
```

```
## # A tibble: 32 x 11  
##   mpg   cyl  disp    hp  drat    wt  qsec    vs    am  
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 21       6   160   110   3.9   2.62  16.5     0     1  
## 2 21       6   160   110   3.9   2.88  17.0     0     1  
## 3 22.8     4   108    93   3.85  2.32  18.6     1     1  
## 4 21.4     6   258   110   3.08  3.22  19.4     1     0  
## 5 18.7     8   360   175   3.15  3.44  17.0     0     0  
## 6 18.1     6   225   105   2.76  3.46  20.2     1     0  
## 7 14.3     8   360   245   3.21  3.57  15.8     0     0  
## 8 24.4     4   147.    62   3.69  3.19   20      1     0  
## 9 22.8     4   141.    95   3.92  3.15  22.9     1     0  
## 10 19.2    6   168.   123   3.92  3.44  18.3     1     0  
## # ... with 22 more rows, and 2 more variables: gear <dbl>,  
## #   carb <dbl>
```



# new data alert!



## mtcars

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1

Where does it come from?

base R

How can I use it?

`View(mtcars)`

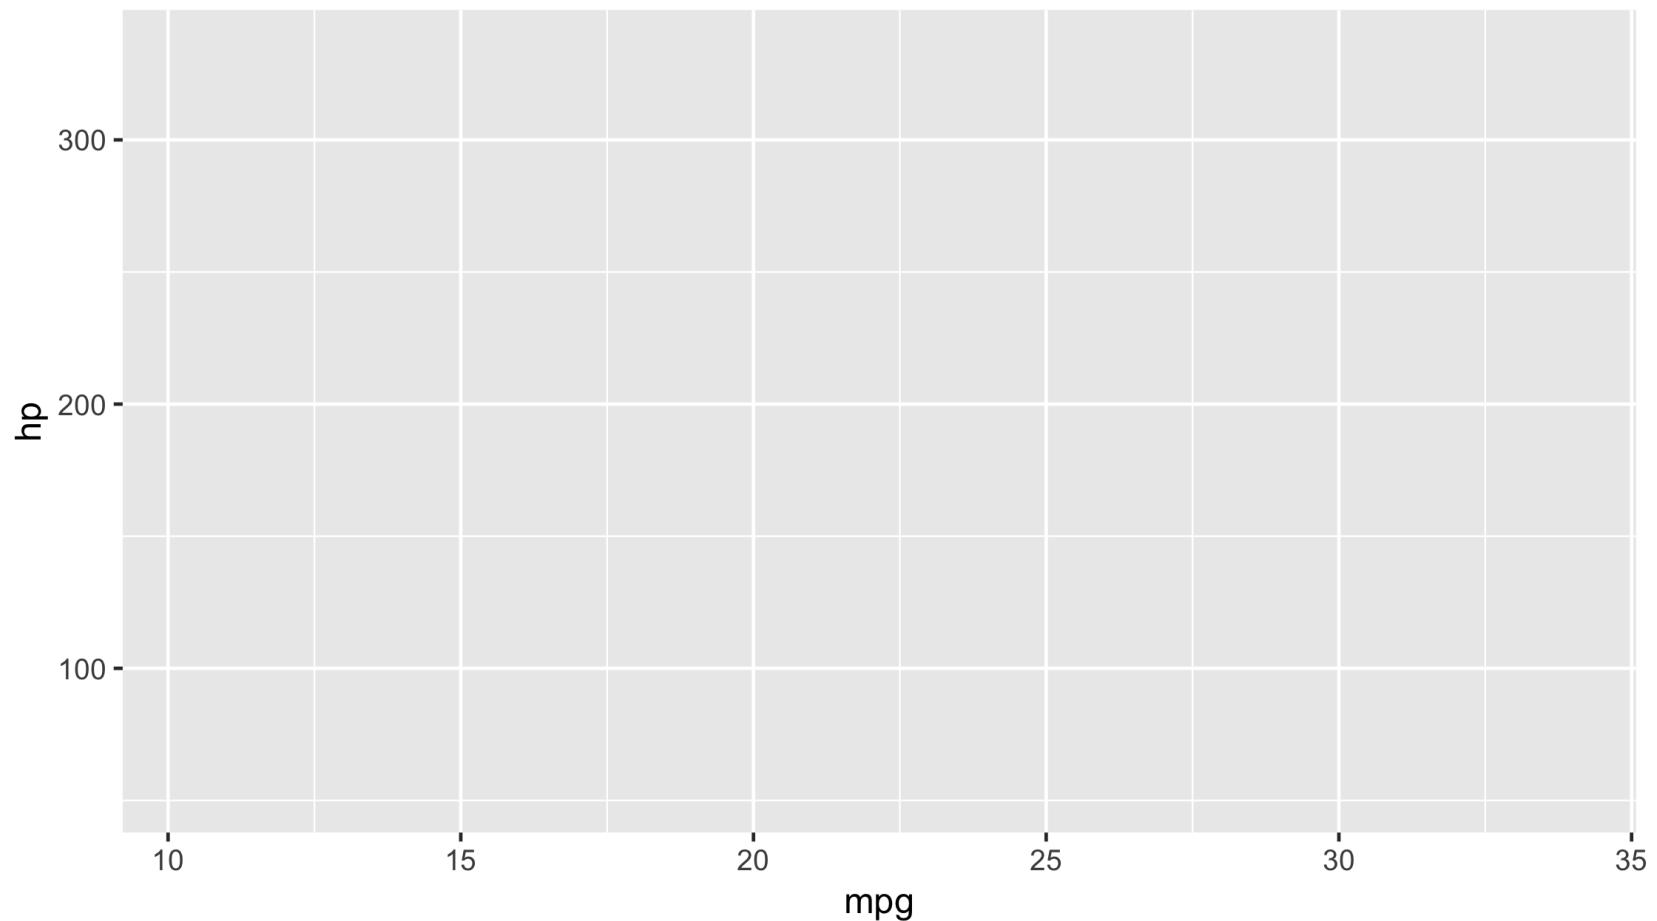


*it's invisible!*

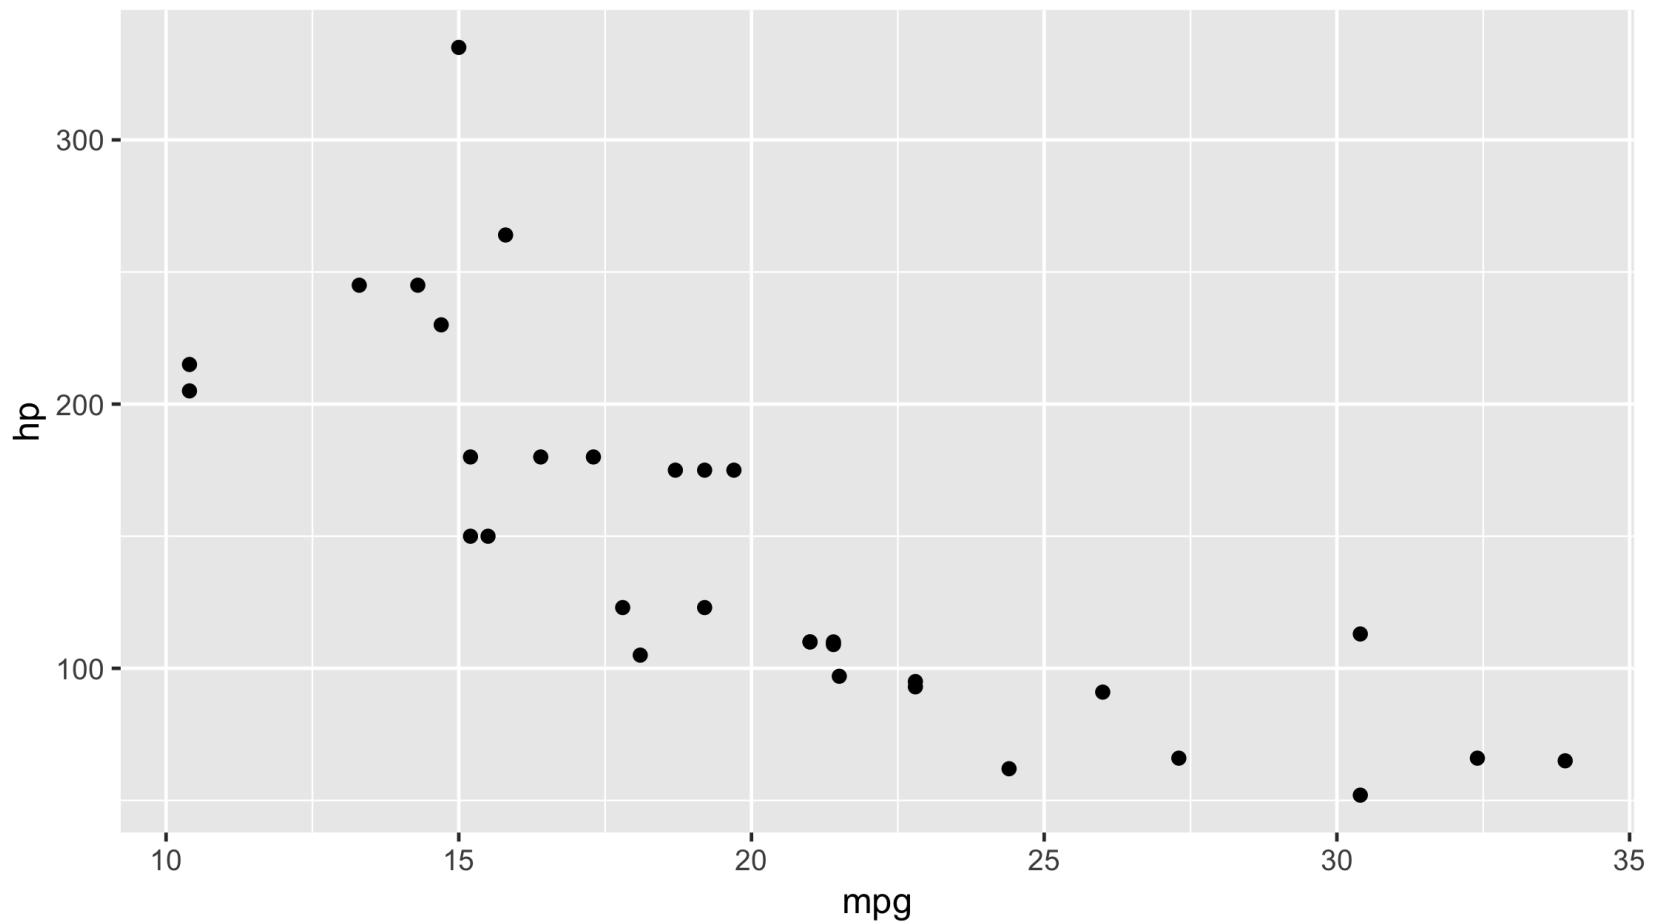
```
ggplot()
```

```
ggplot()
```

```
ggplot(mtcars, aes(x = mpg, y = hp))
```



```
ggplot(mtcars, aes(x = mpg, y = hp)) +  
  geom_point()
```



# ggplot()

```
ggplot(data = <data>, mapping = aes(<mapping>)) +  
<geom_function>()
```

# ggplot()

```
ggplot(data = <data>, mapping = aes(<mapping>)) +  
<geom_function>()
```

Add layers with +

# ggplot()

```
ggplot(data = <data>, mapping = aes(<mapping>)) +  
<geom_function>()
```

Add layers with +

Put + at the **end** of a line

# ggplot()

```
ggplot(data = <data>, mapping = aes(<mapping>)) +  
<geom_function>()
```

Add layers with +

Put + at the end of a line

map aesthetics with aes()

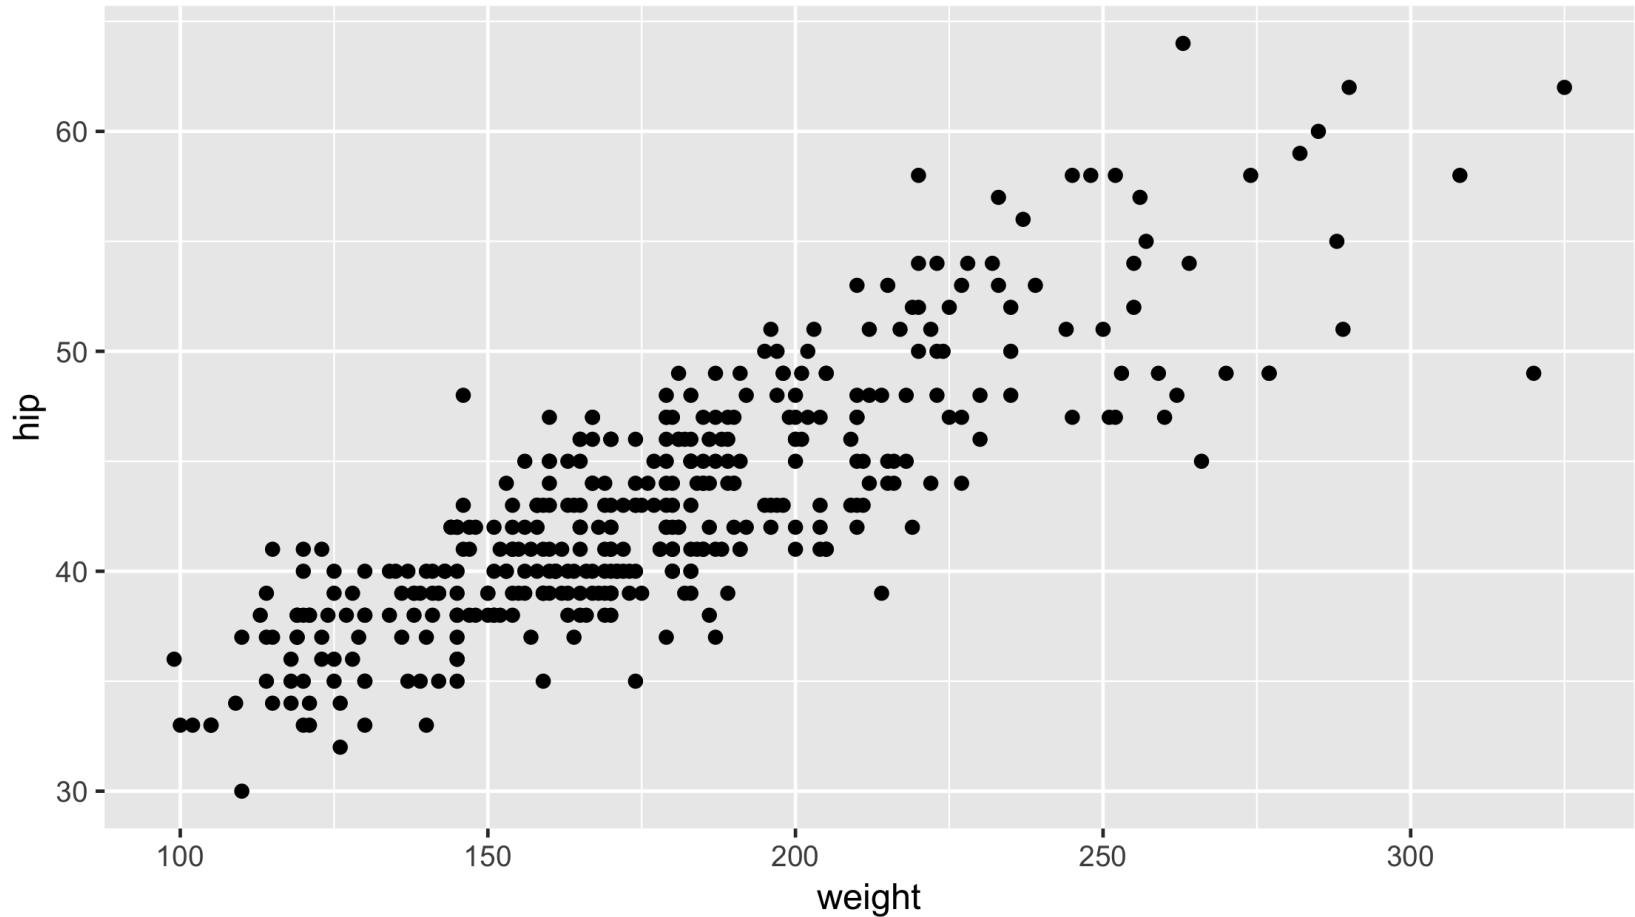
# Your Turn 1

Read in the diabetes data.

Write and run the code from this slide to make a graph. Pay strict attention to spelling, capitalization, and parentheses!

```
ggplot(data = diabetes, mapping = aes(x = weight, y = hip)) +  
  geom_point()
```

```
diabetes <- read_csv("diabetes.csv")  
  
ggplot(data = diabetes, mapping = aes(x = weight, y = hip)) +  
  geom_point()
```



# Aesthetics: aes()

```
ggplot(data = <data>, mapping = aes(<mapping>)) +  
<geom_function>()
```

Aesthetics **map** the data to the plot

# Aesthetics: aes()

```
ggplot(mtcars, aes(x = mpg, y = hp, color = cyl)) + geom_point()  
ggplot(mtcars, aes(x = mpg, y = hp, size = cyl)) + geom_point()  
ggplot(mtcars, aes(x = mpg, y = hp, alpha = cyl)) + geom_point()  
ggplot(mtcars, aes(x = mpg, y = hp, shape = cyl)) + geom_point()
```

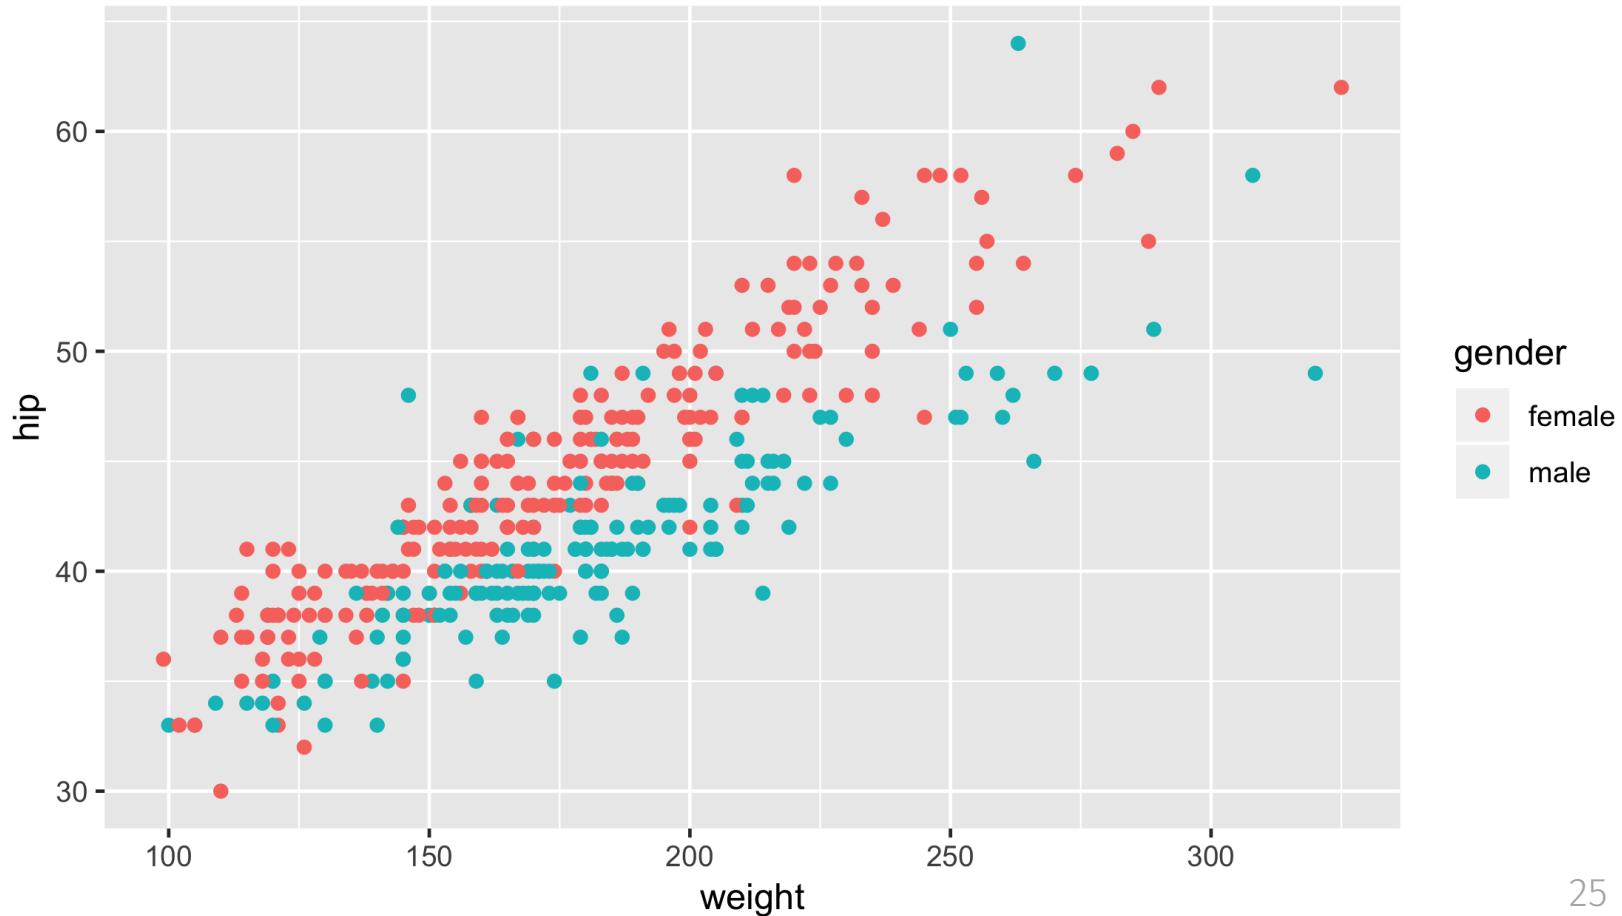
## Your Turn 2

Add color, size, alpha, and shape aesthetics to your graph using the gender variable. Experiment.

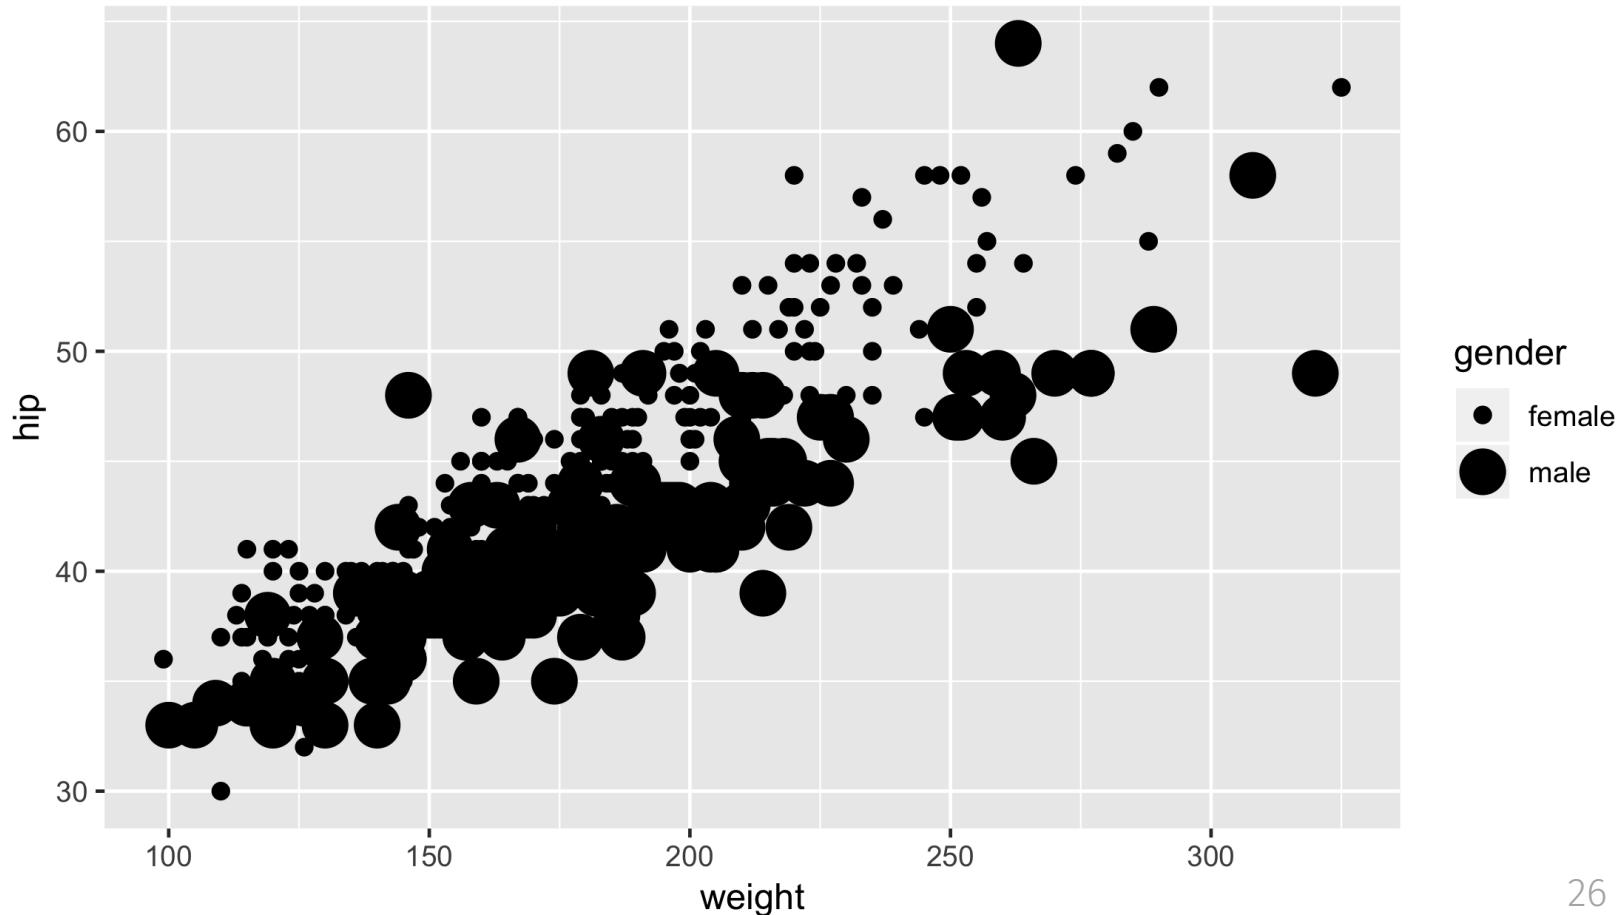
```
ggplot(data = diabetes,  
       mapping = aes(x = weight, y = hip)) +  
       geom_point()
```

Try moving the mapping argument to `geom_point()`. Add in any aesthetics you found helpful.

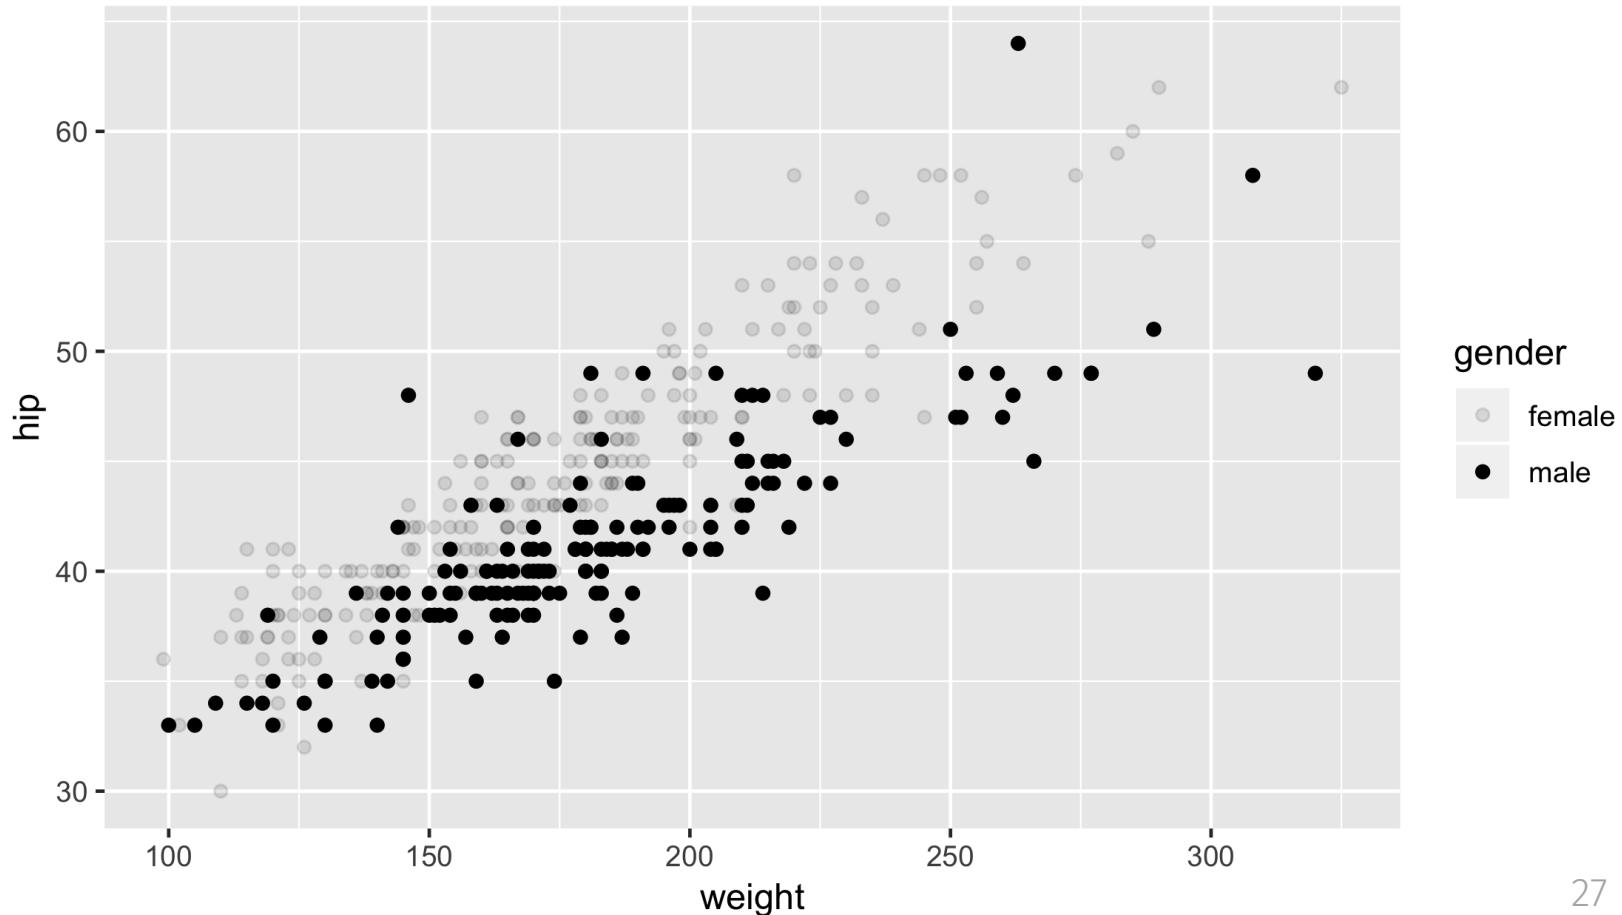
```
ggplot(  
  data = diabetes,  
  mapping = aes(x = weight, y = hip, color = gender)  
) +  
  geom_point()
```



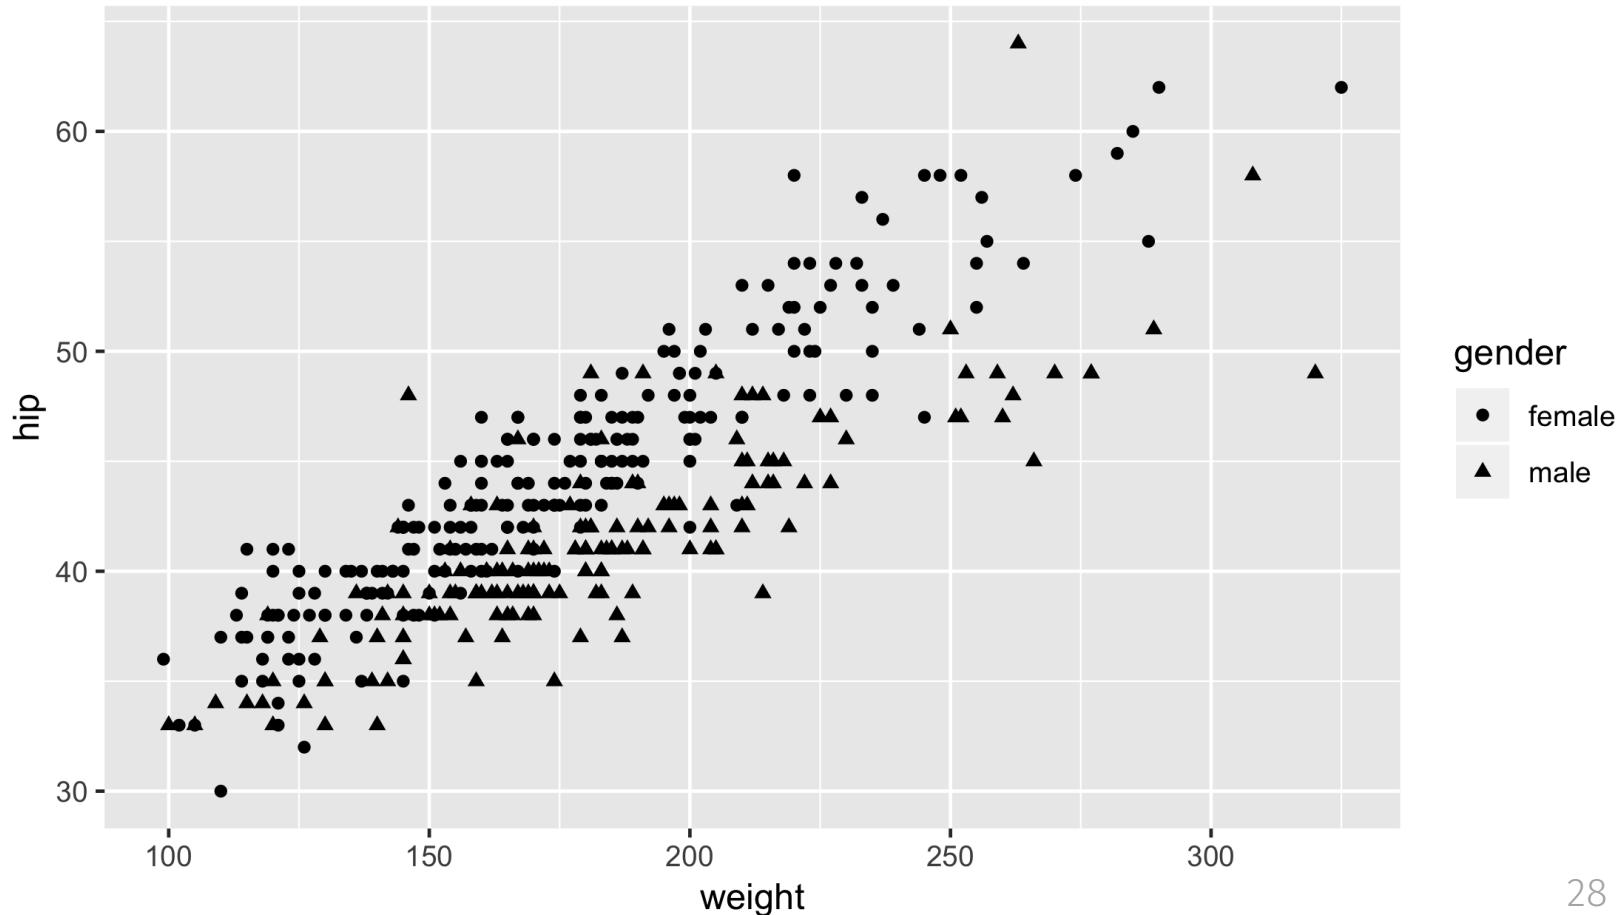
```
ggplot(  
  data = diabetes,  
  mapping = aes(x = weight, y = hip, size = gender)  
) +  
  geom_point()
```



```
ggplot(  
  data = diabetes,  
  mapping = aes(x = weight, y = hip, alpha = gender)  
) +  
  geom_point()
```



```
ggplot(  
  data = diabetes,  
  mapping = aes(x = weight, y = hip, shape = gender)  
) +  
  geom_point()
```



```
ggplot(data = diabetes) +  
  geom_point(  
    mapping = aes(  
      x = weight,  
      y = hip,  
      color = gender,  
      shape = gender  
    )  
  )
```

# geoms

What shape does the data take?

geom\_point()

geom\_line()

geom\_violin()

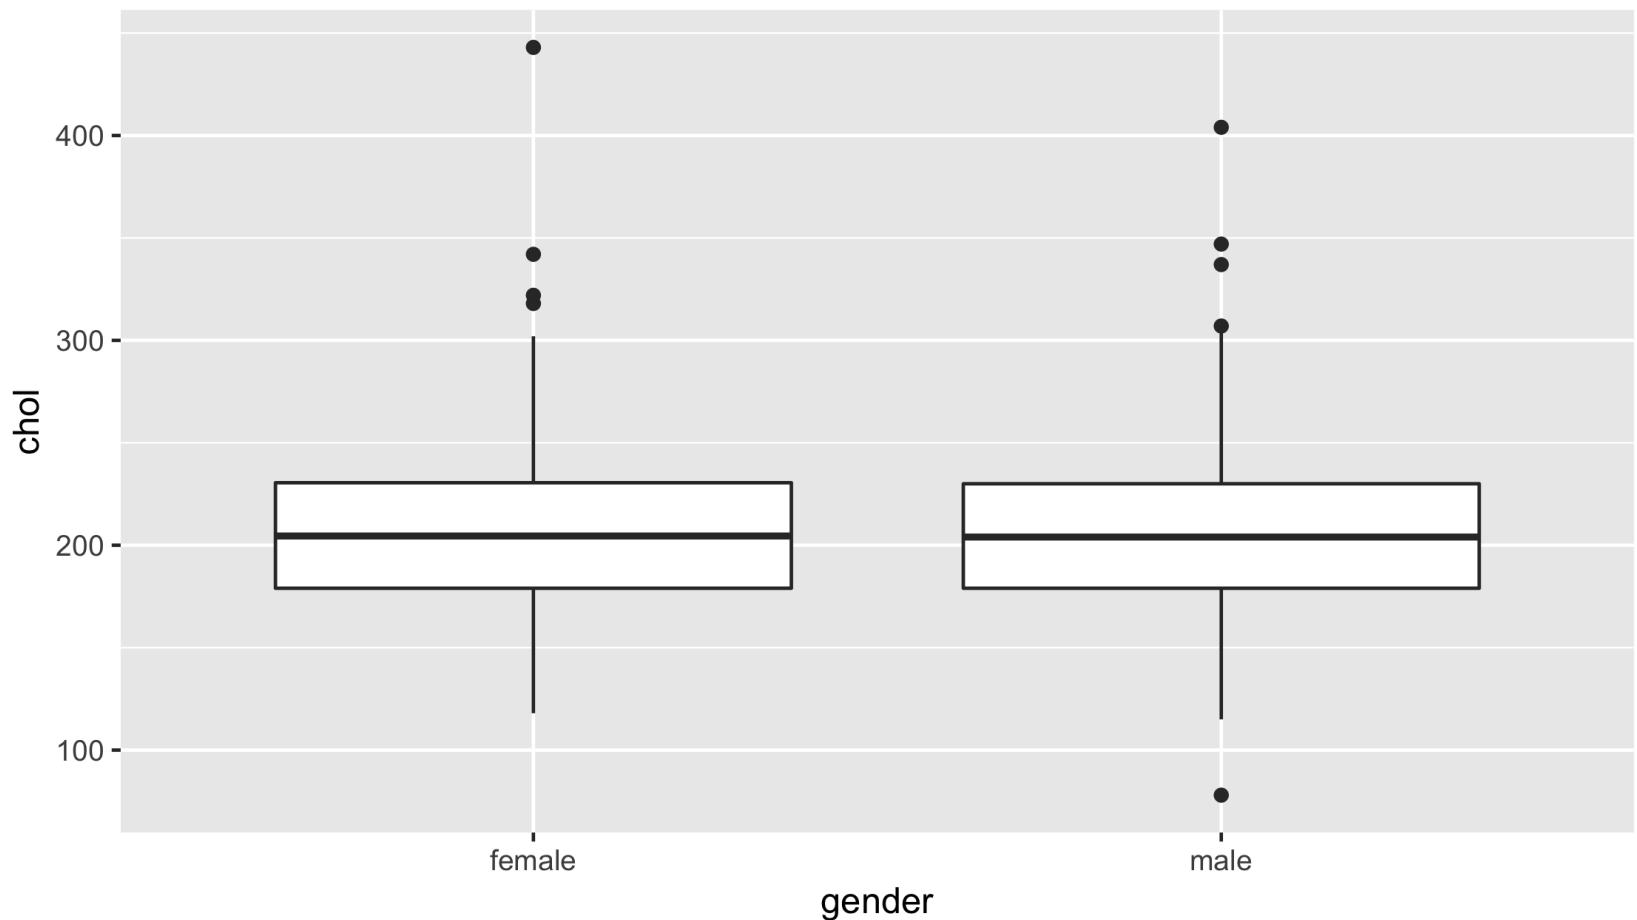
**Check the cheatsheet!**

# Your Turn 3

Replace this scatterplot with one that draws boxplots.

```
ggplot(diabetes, aes(gender, chol)) + geom_point()
```

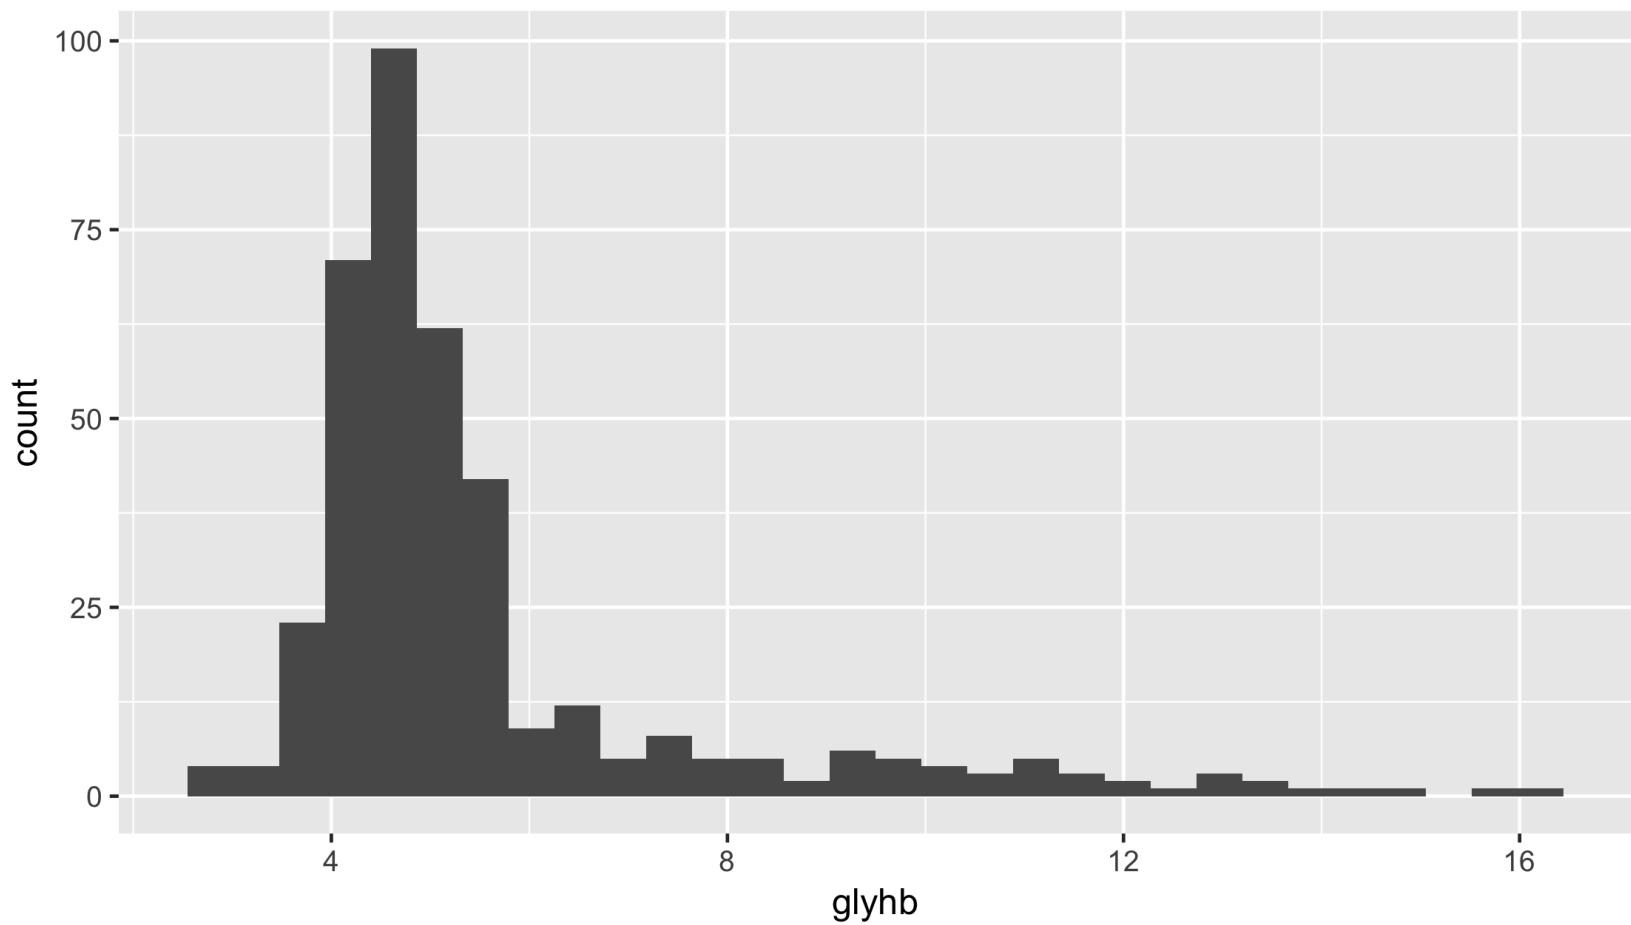
```
ggplot(diabetes, aes(gender, chol)) + geom_boxplot()
```



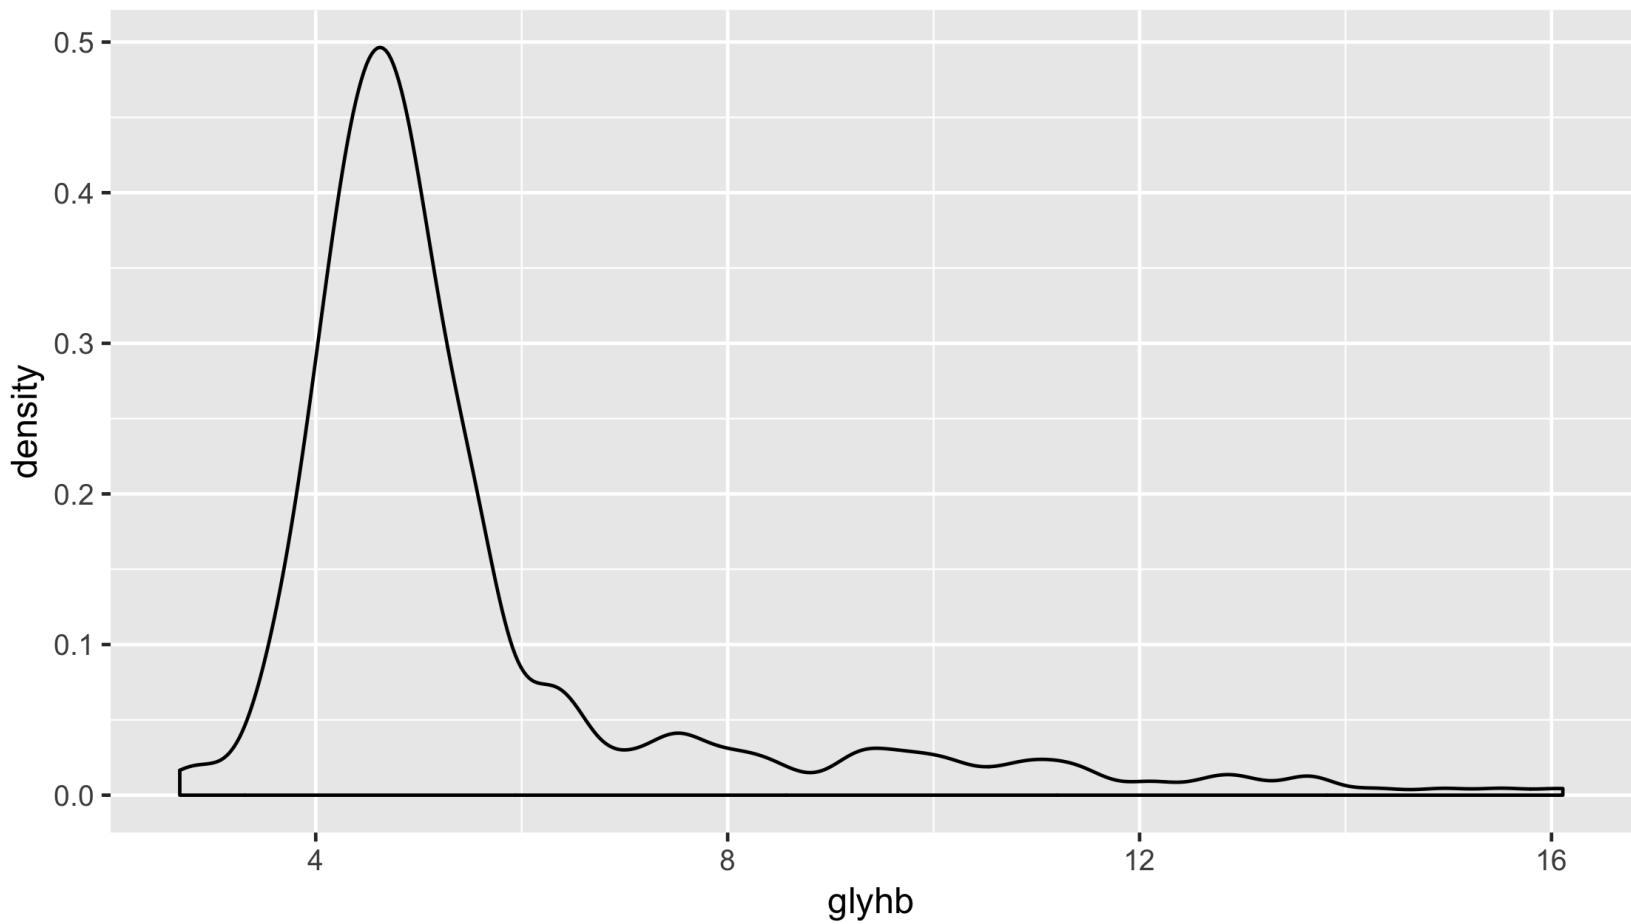
## Your Turn 4

1. Make a histogram of the glyhb variable in diabetes.
2. Redo the glyhb plot as a density plot.

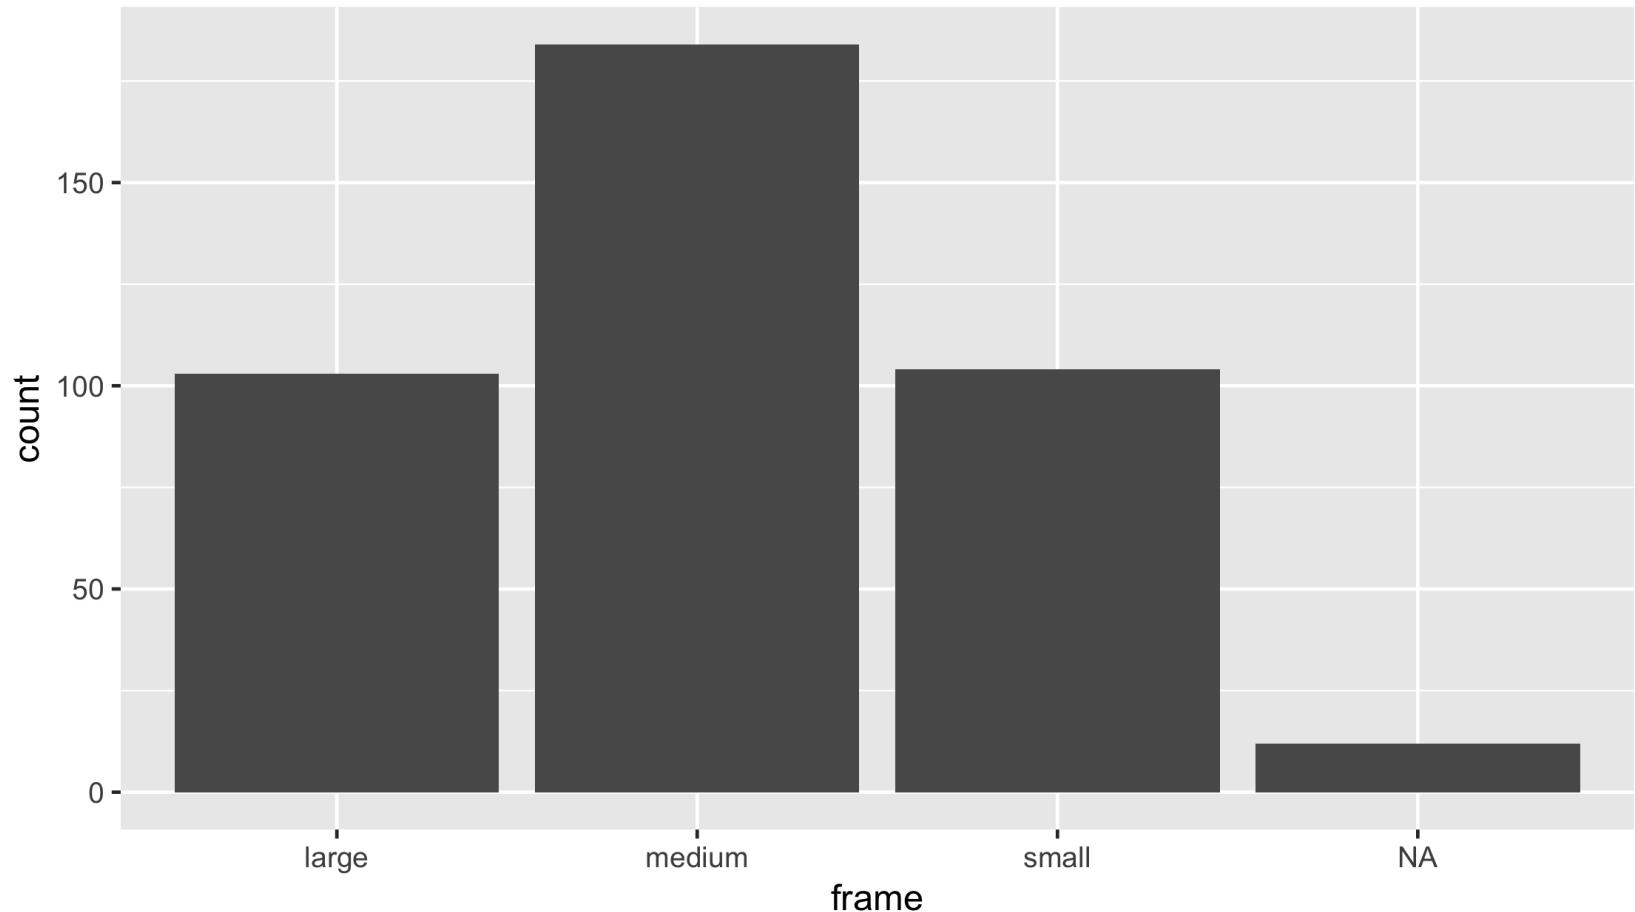
```
ggplot(diabetes, aes(x = glyhb)) +  
  geom_histogram()
```



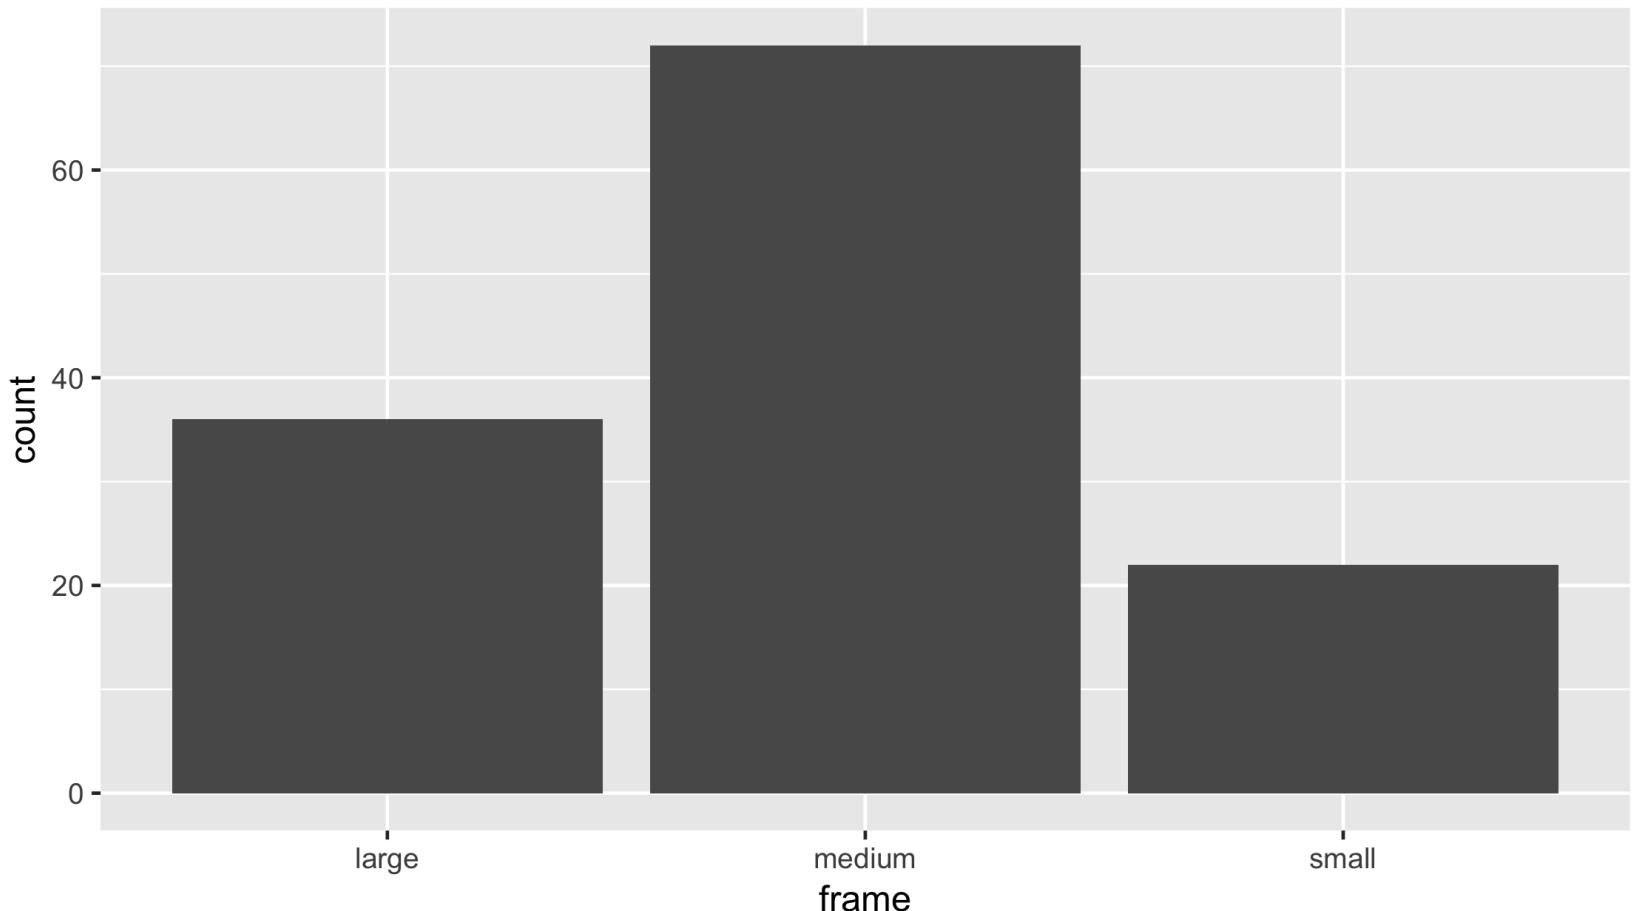
```
ggplot(diabetes, aes(x = glyhb)) +  
  geom_density()
```



```
diabetes %>%
  ggplot(aes(x = frame)) +
  geom_bar()
```



```
diabetes %>%
  drop_na() %>%
  ggplot(aes(x = frame)) +
  geom_bar()
```

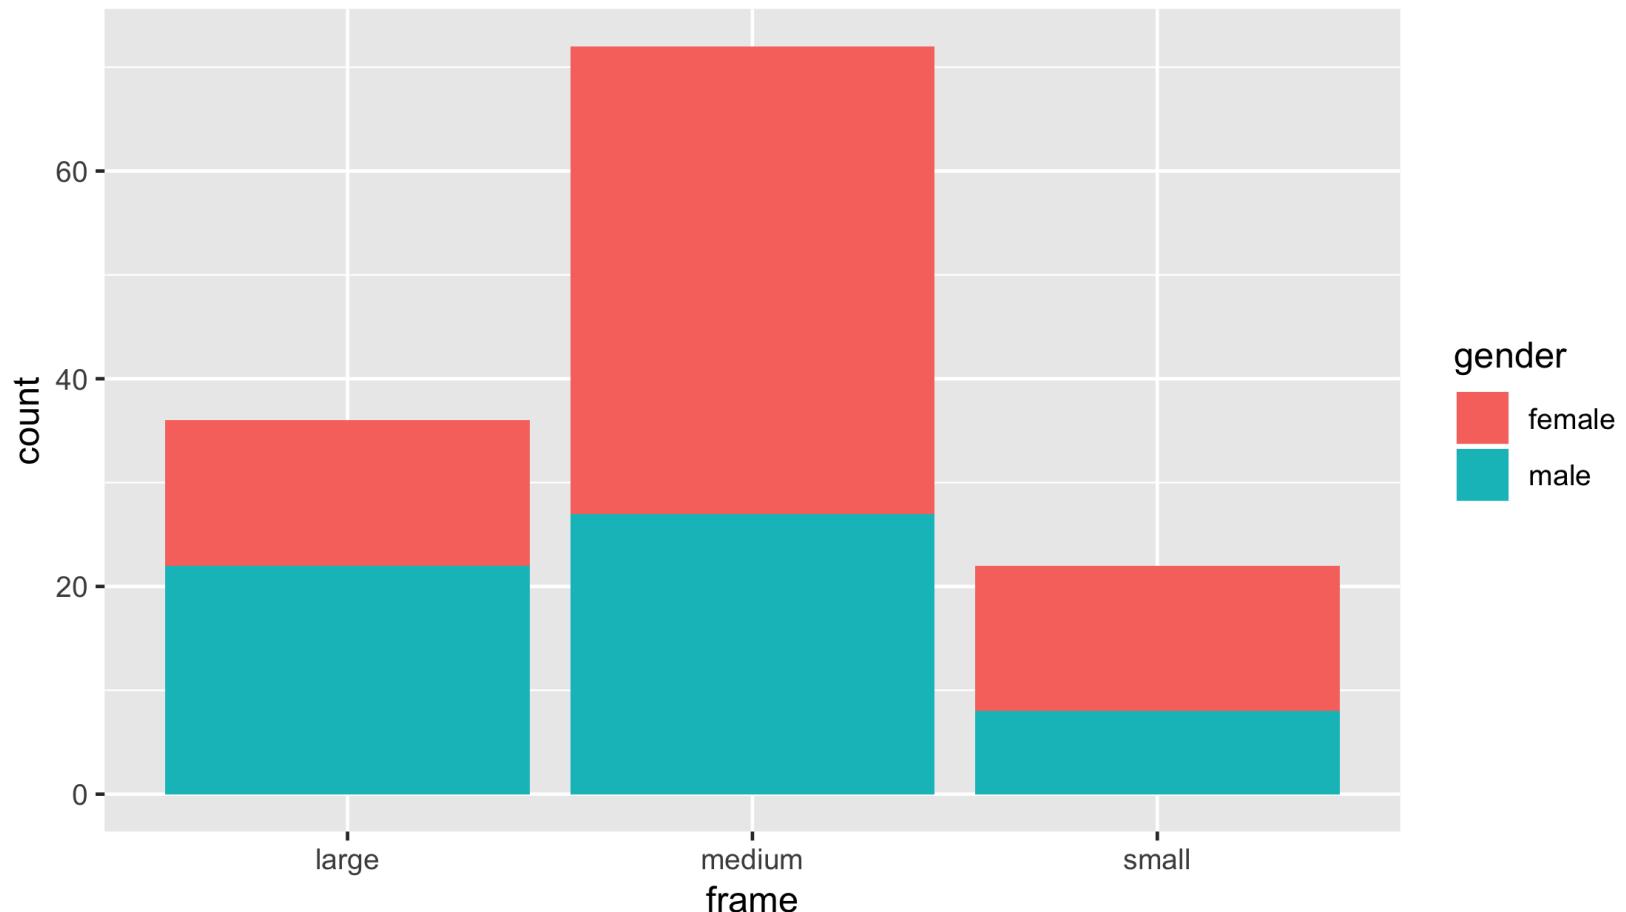


## Your Turn 5

Make a bar chart of frame colored by gender. Then, try it with the fill aesthetic instead of color.

```
diabetes %>%
  drop_na() %>%
  -----() +
  -----()
```

```
diabetes %>%
  drop_na() %>%
  ggplot(aes(x = frame, fill = gender)) +
  geom_bar()
```



# Positions

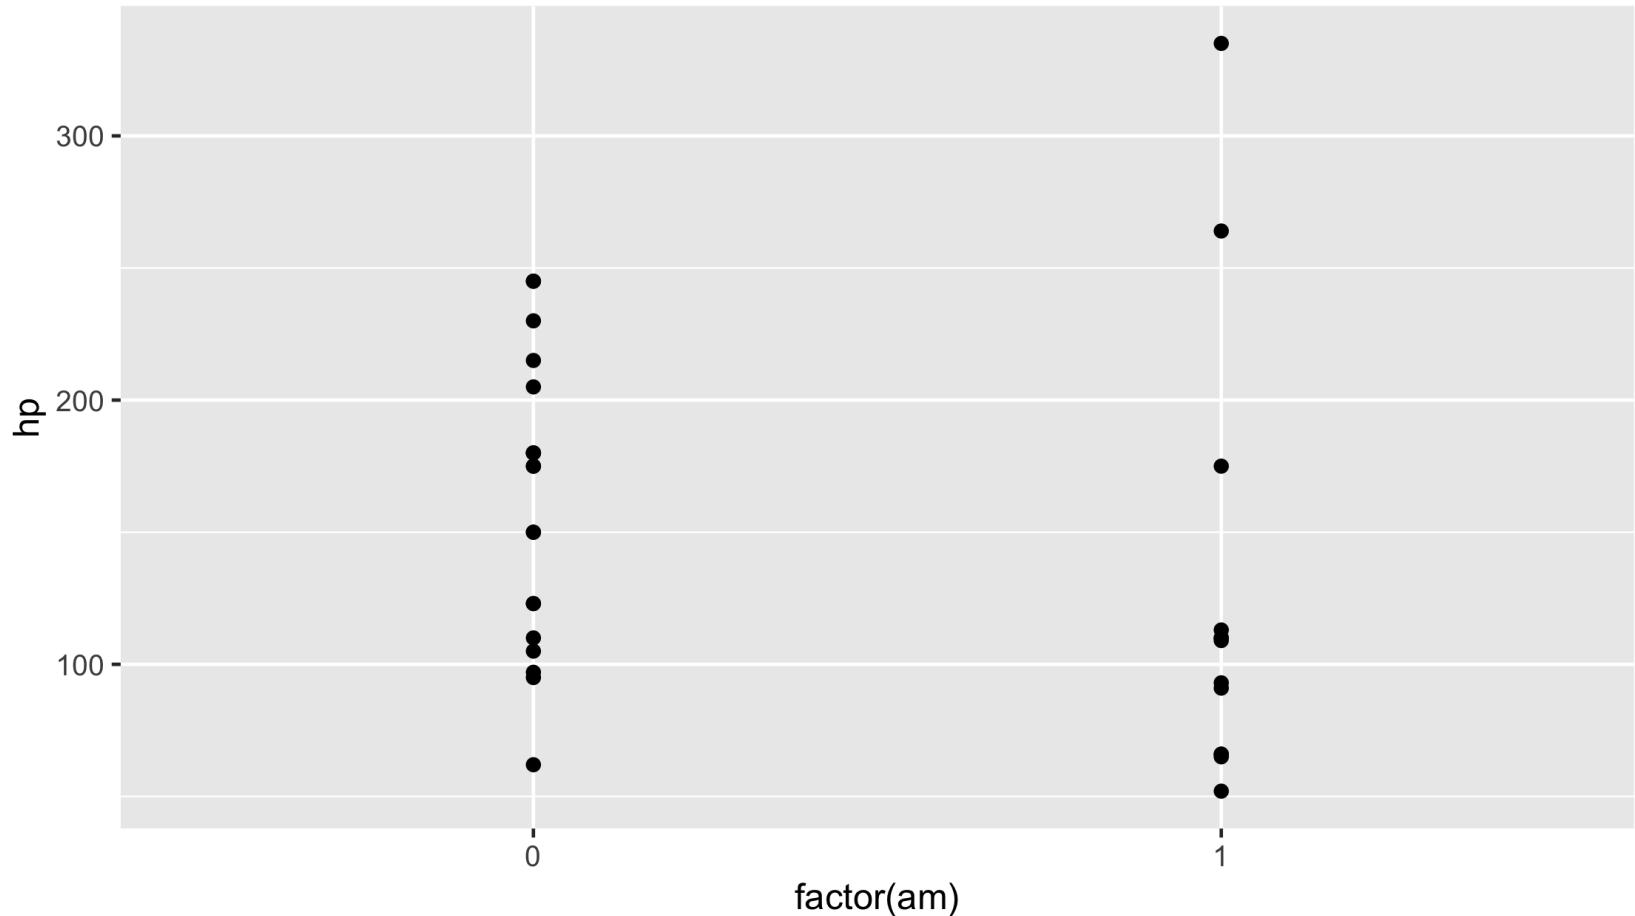
`geom_bar(position = "<POSITION>")`

When we have aesthetics mapped, how are they positioned?

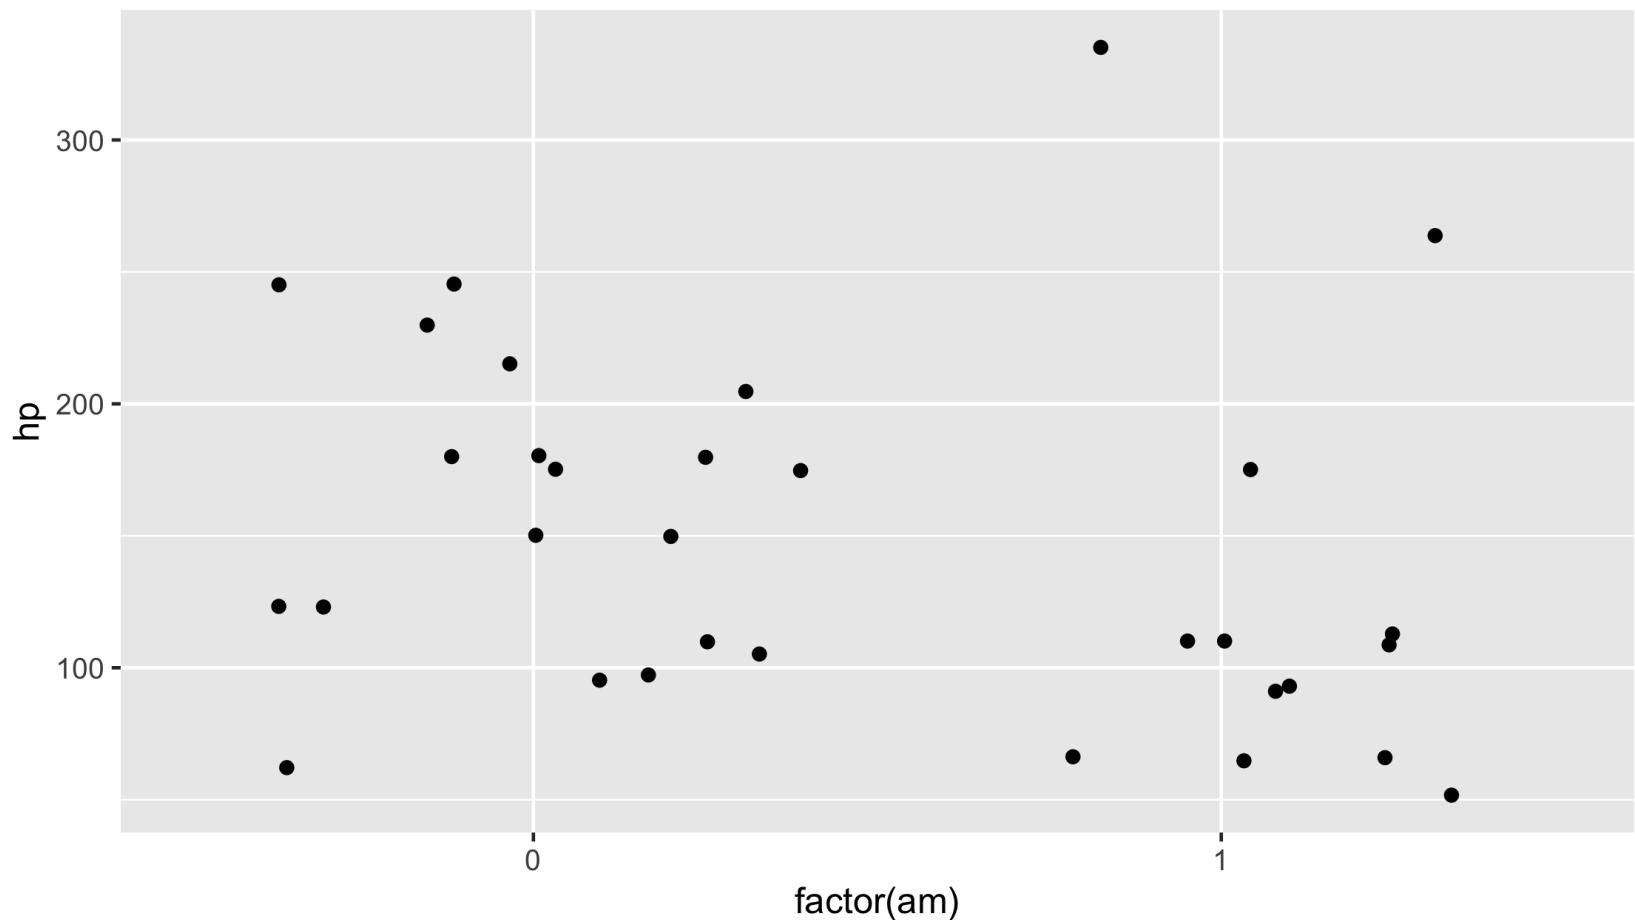
**geom\_bar: dodge, fill, stacked (default)**

**geom\_point: jitter**

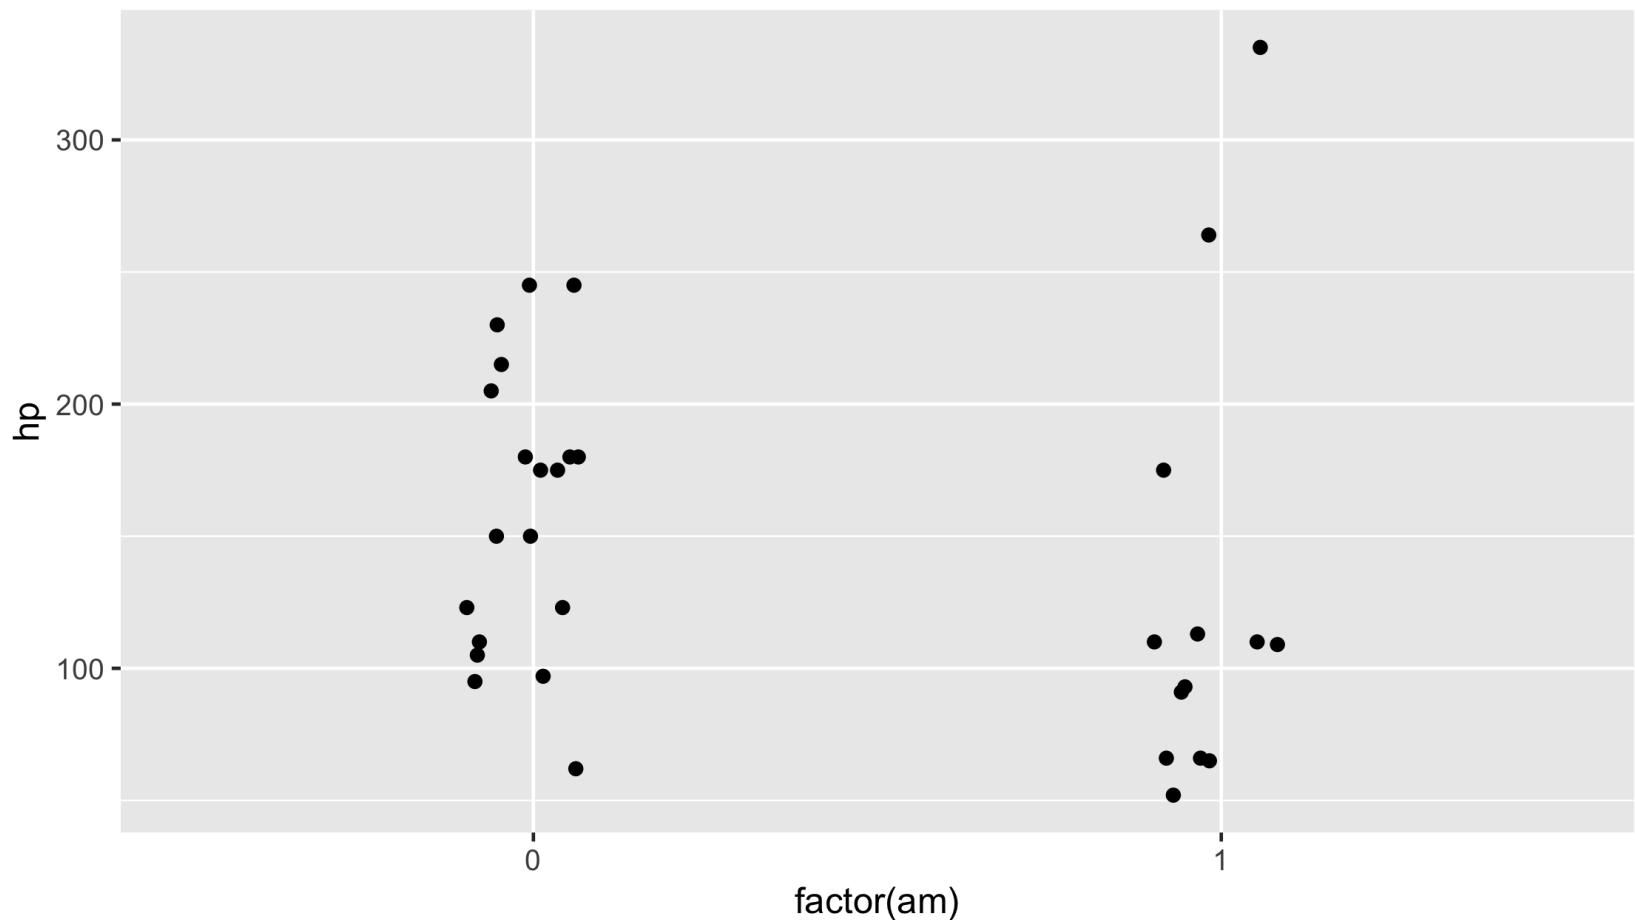
```
ggplot(mtcars, aes(x = factor(am), y = hp)) +  
  geom_point()
```



```
ggplot(mtcars, aes(x = factor(am), y = hp)) +  
  geom_point(position = "jitter")
```



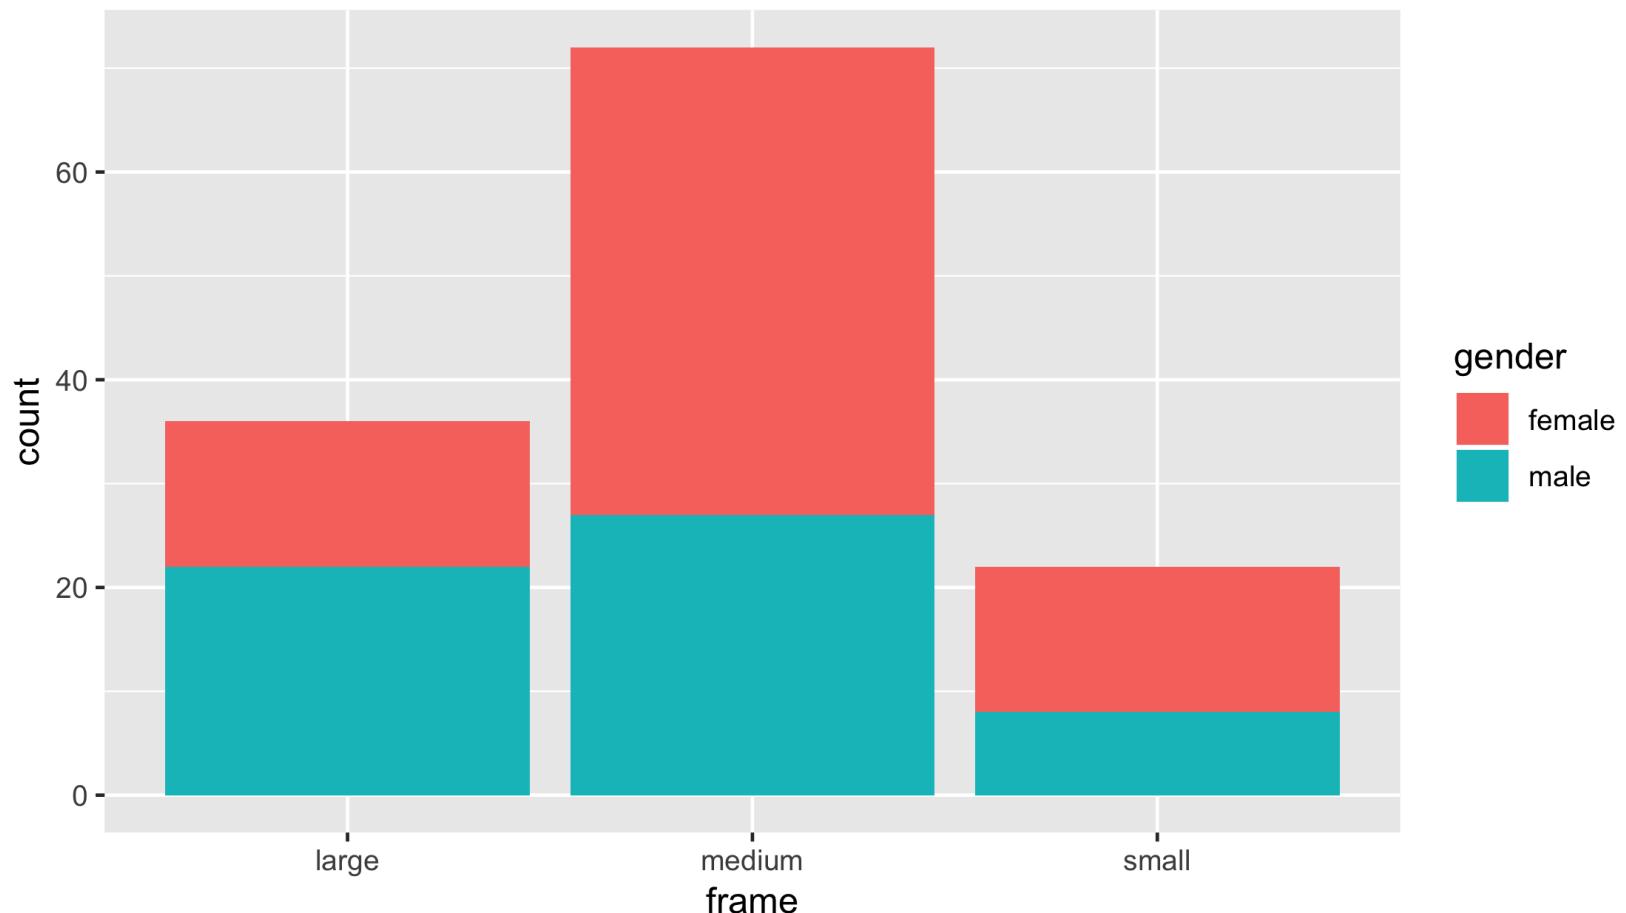
```
ggplot(mtcars, aes(x = factor(am), y = hp)) +  
  geom_jitter(width = .1, height = 0)
```



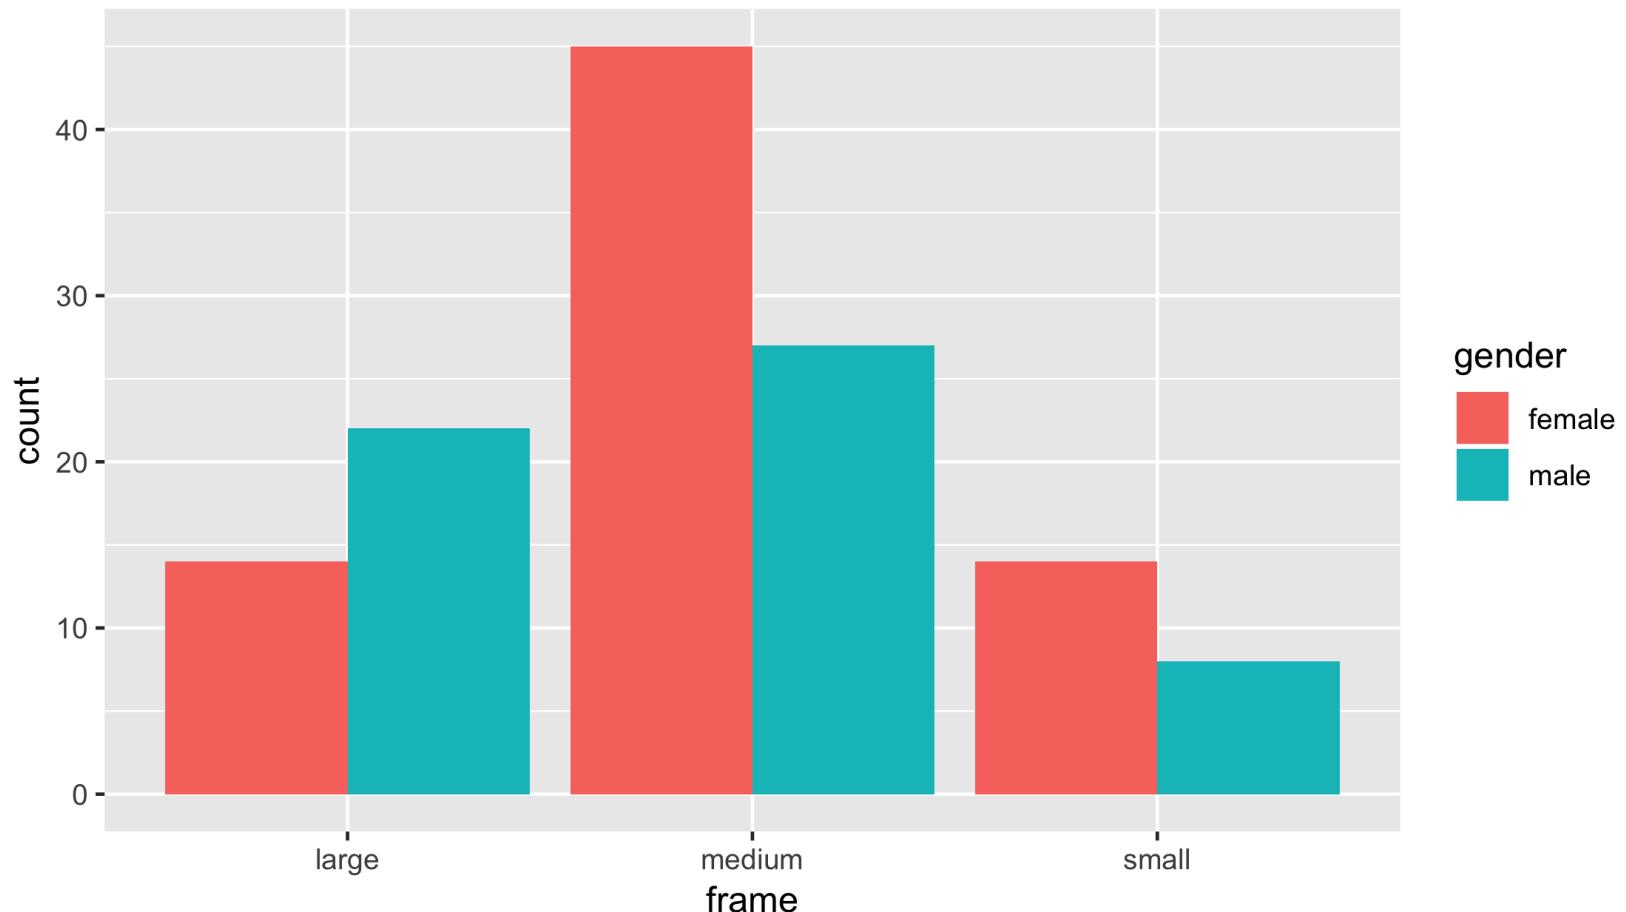
## Your Turn 6

Take your code for the bar chart before (using the fill aesthetic). Experiment with different position values: "dodge", "fill", "stack"

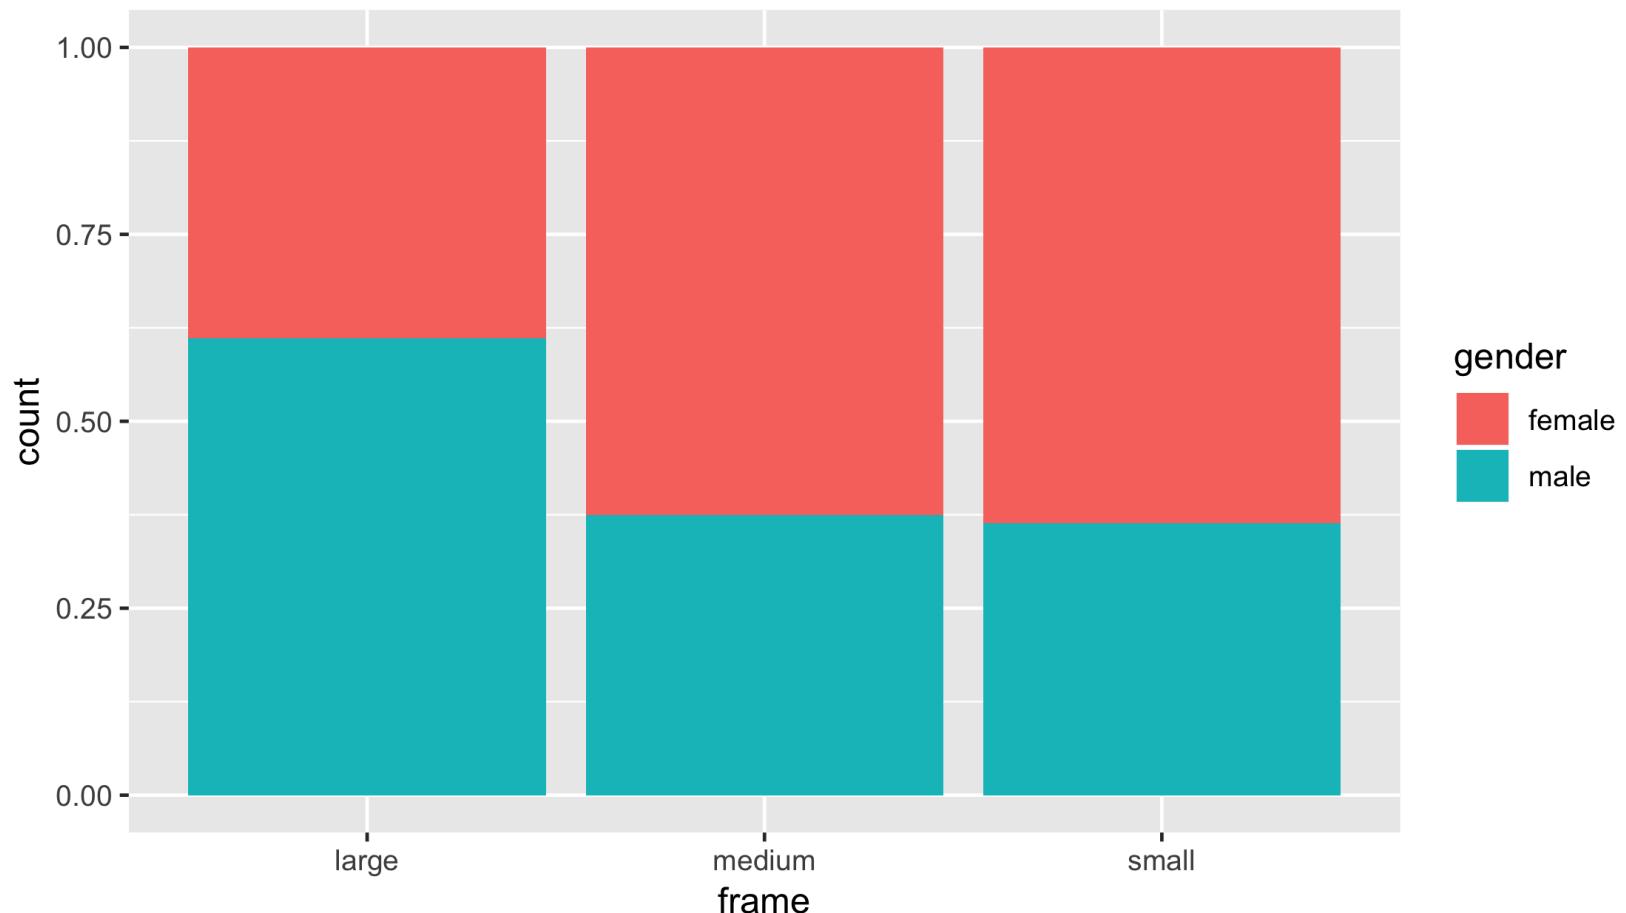
```
diabetes %>%
  drop_na() %>%
  ggplot(aes(x = frame, fill = gender)) +
  geom_bar(position = "stack")
```



```
diabetes %>%
  drop_na() %>%
  ggplot(aes(x = frame, fill = gender)) +
  geom_bar(position = "dodge")
```



```
diabetes %>%
  drop_na() %>%
  ggplot(aes(x = frame, fill = gender)) +
  geom_bar(position = "fill")
```



# Mapping vs setting

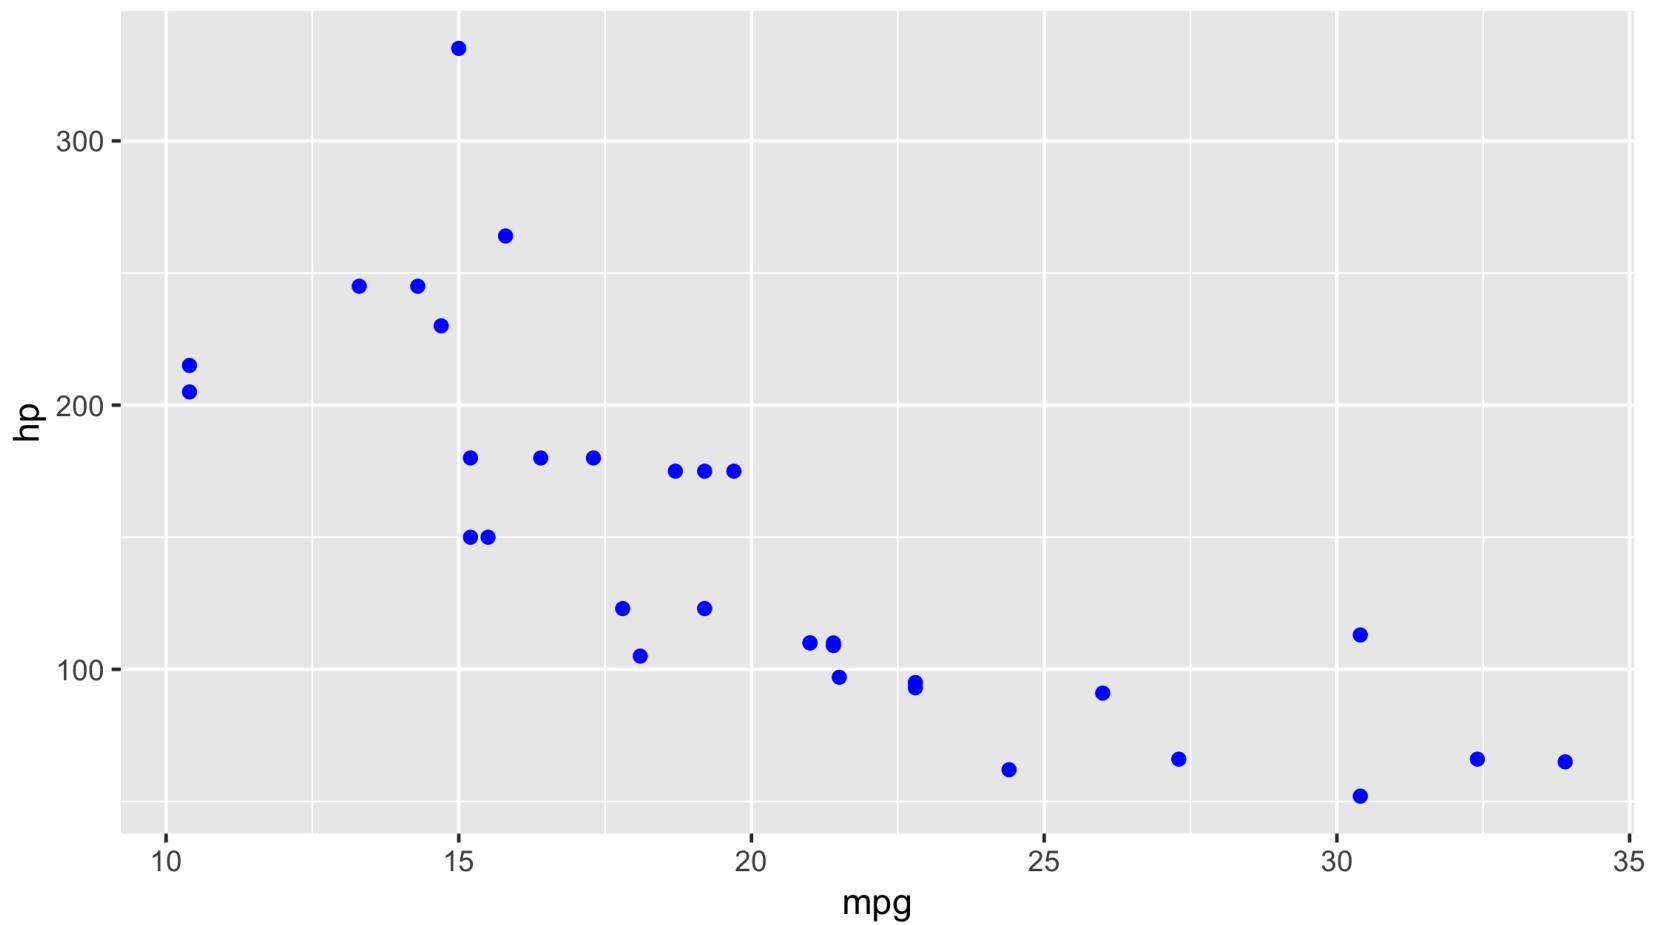
Cool, but how do I just make everything  
**blue?**

```
geom_point(aes(x, y), color = "blue")
```

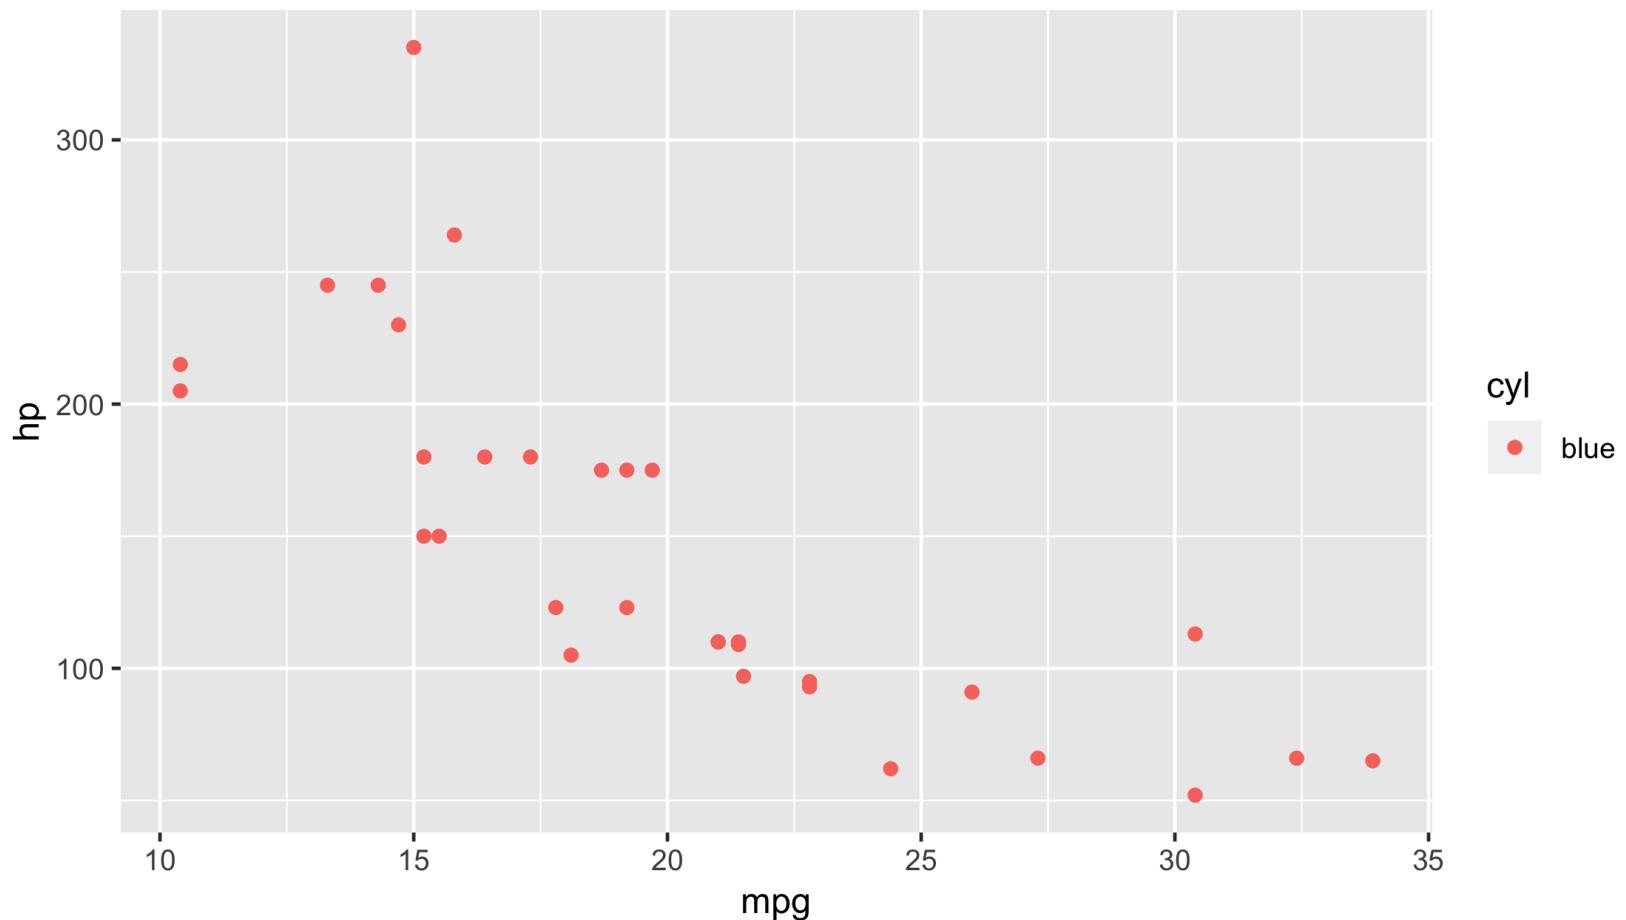
```
geom_bar(aes(x, y), fill = "blue")
```

To set a color, put it **outside** aes()

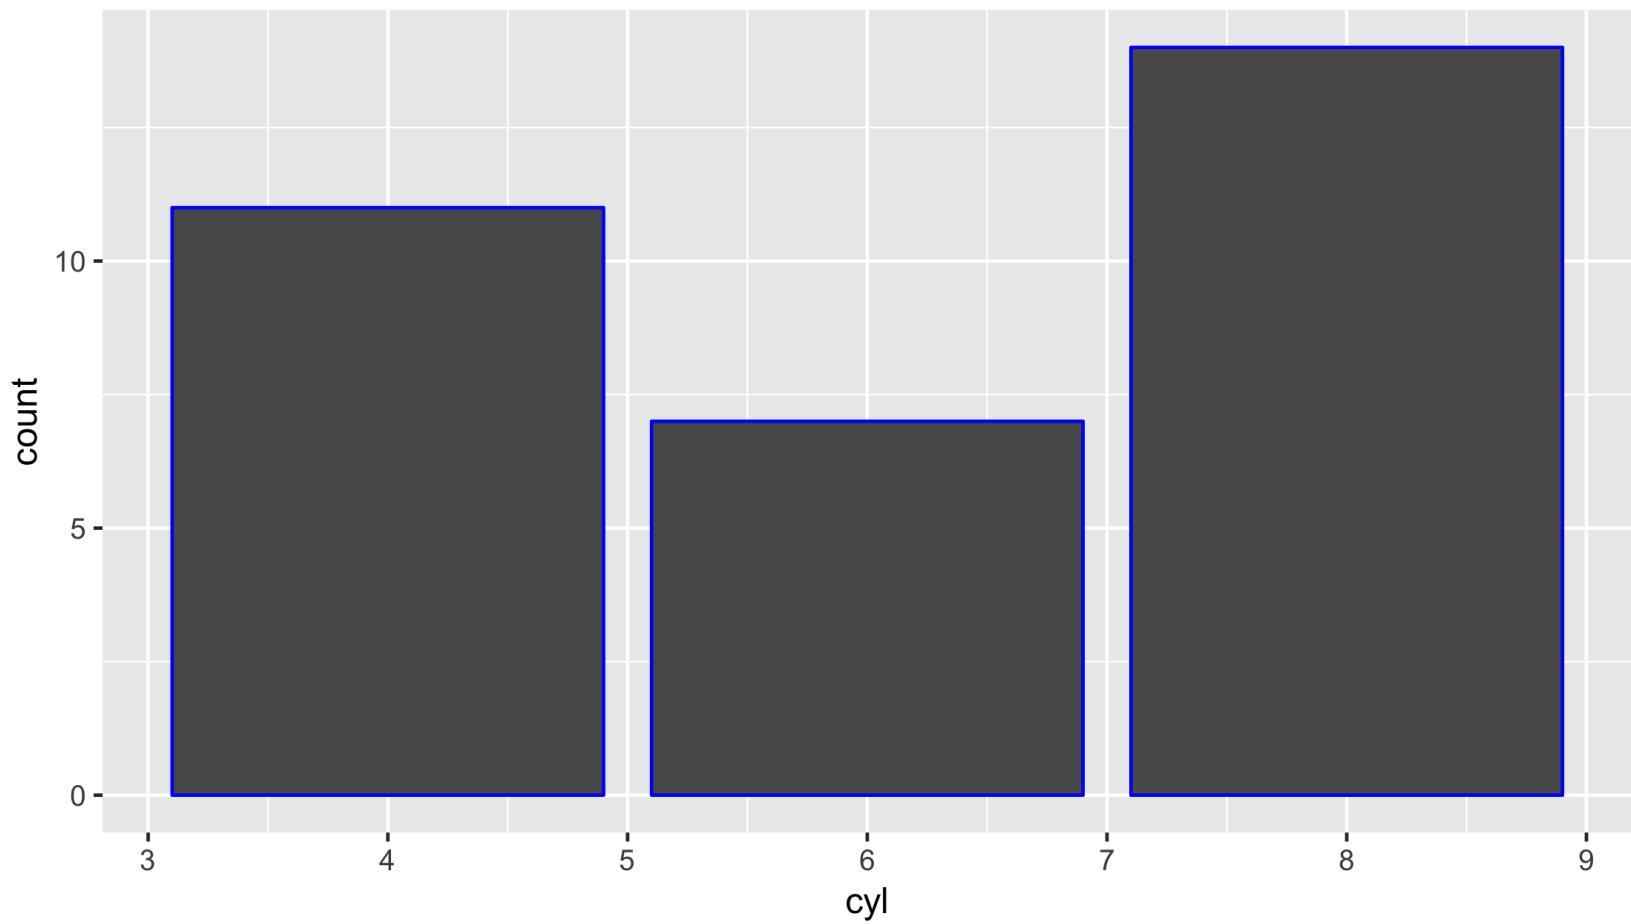
```
ggplot(mtcars, aes(x = mpg, y = hp, color = cyl)) +  
  geom_point(color = "blue")
```



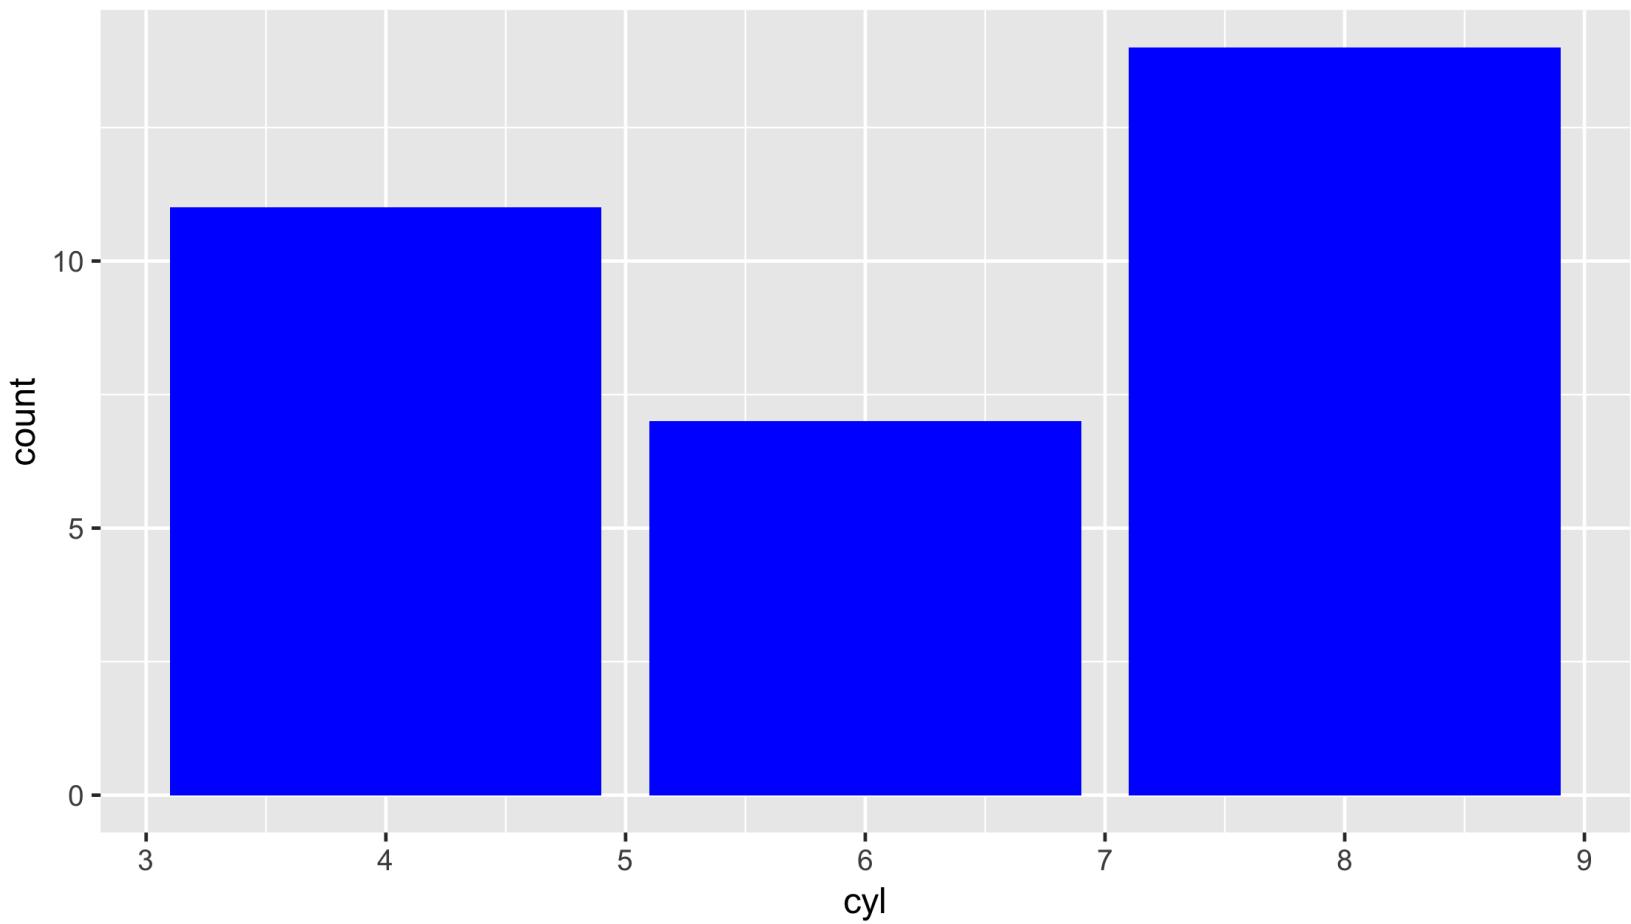
```
ggplot(mtcars, aes(x = mpg, y = hp, color = cyl)) +  
  geom_point(aes(color = "blue"))
```



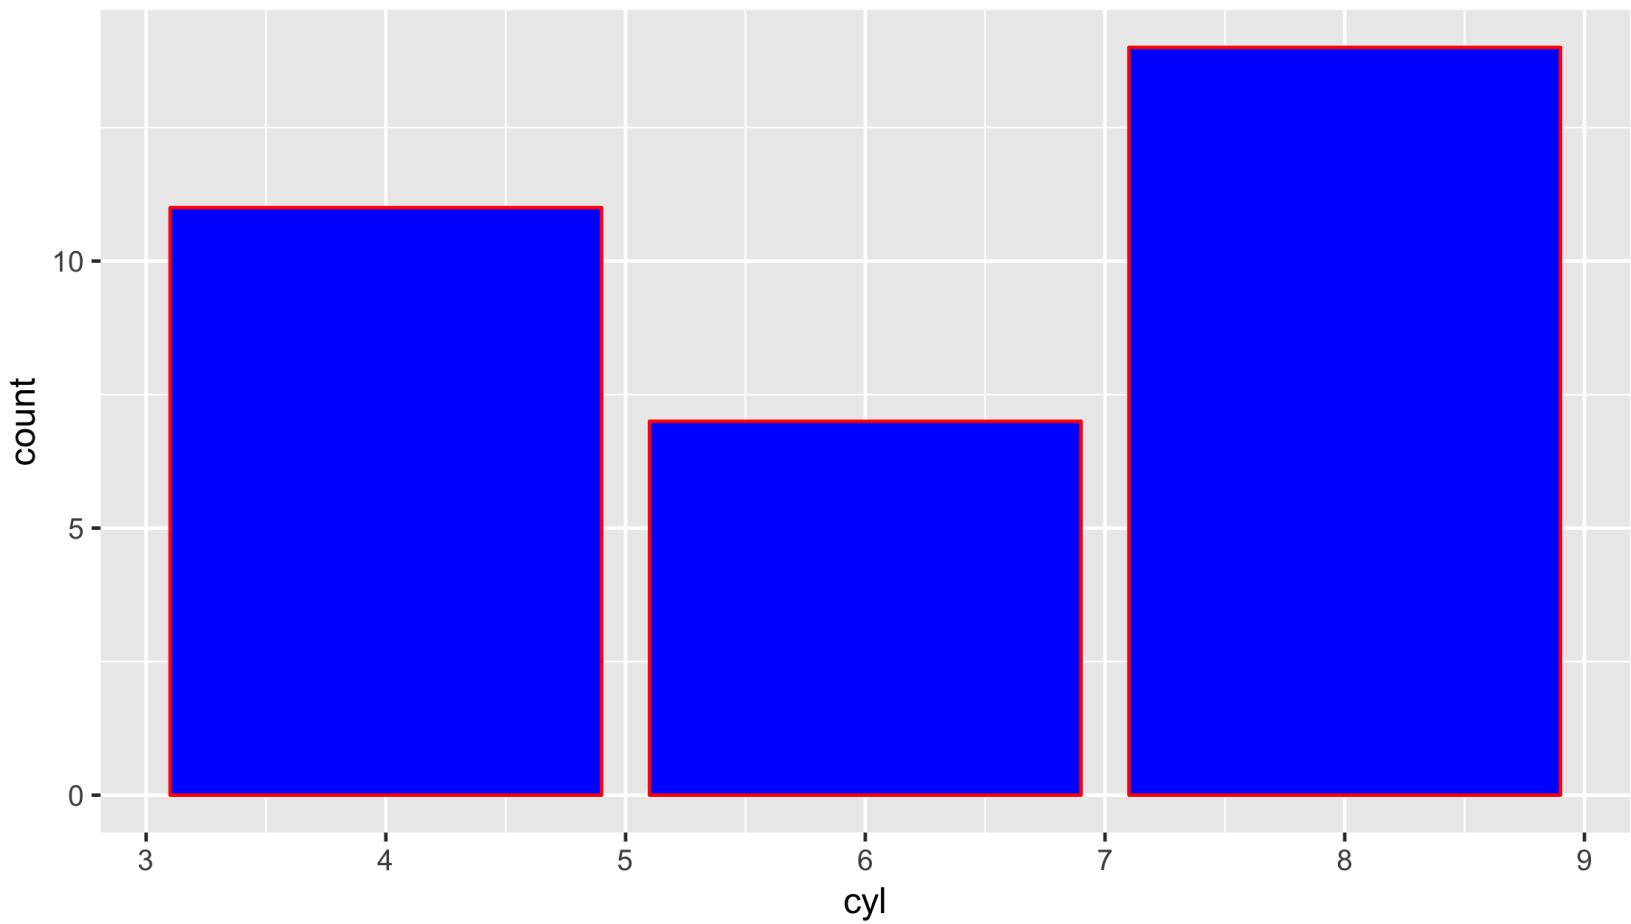
```
ggplot(mtcars, aes(x = cyl)) +  
  geom_bar(color = "blue")
```



```
ggplot(mtcars, aes(x = cyl)) +  
  geom_bar(fill = "blue")
```



```
ggplot(mtcars, aes(x = cyl)) +  
  geom_bar(color = "red", fill = "blue")
```



# Adding layers

```
ggplot(data = <DATA>, mapping = aes(<MAPPINGS>)) +  
<GEOM_FUNCTION>() +  
<GEOM_FUNCTION>() +  
<SCALE_FUNCTION>() +  
<THEME_FUNCTION>()
```

# Your Turn 7

Run the code after every change you make.

1. Predict what this code will do. Then run it.

2. Add a linetype aesthetic for gender. Run it again.

3. Set the color of geom\_smooth() to "black"

4. Add se = FALSE to the geom\_smooth()

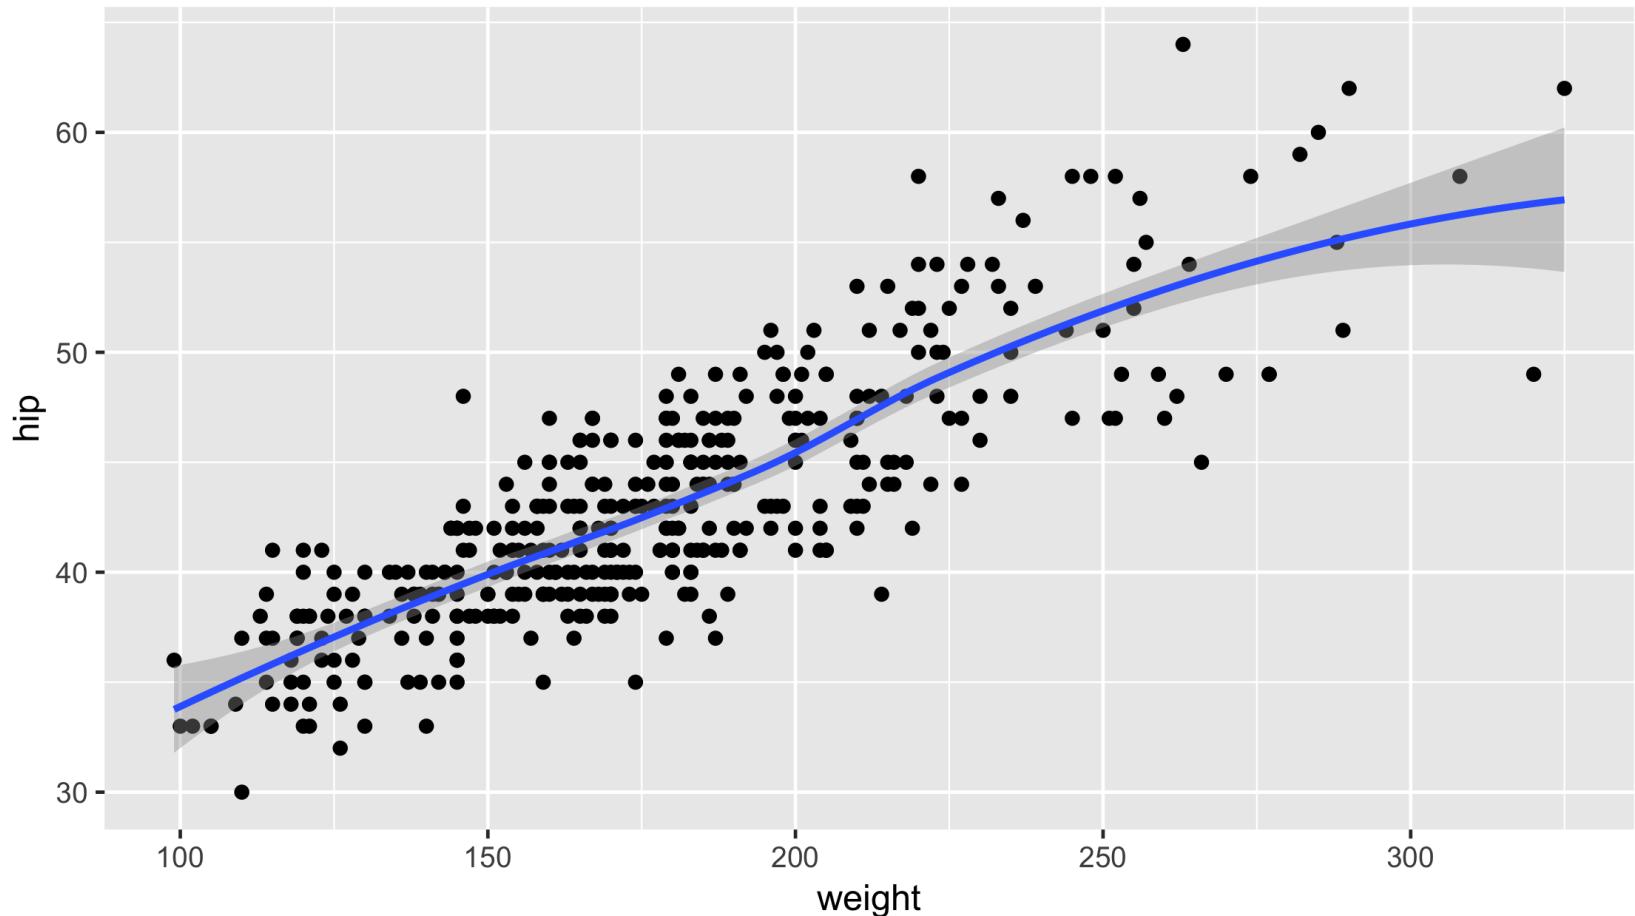
5. It's hard to see the lines well now. How about setting alpha = .2 in geom\_point()?

6. Jitter the points. You can either change the geom or change the position argument.

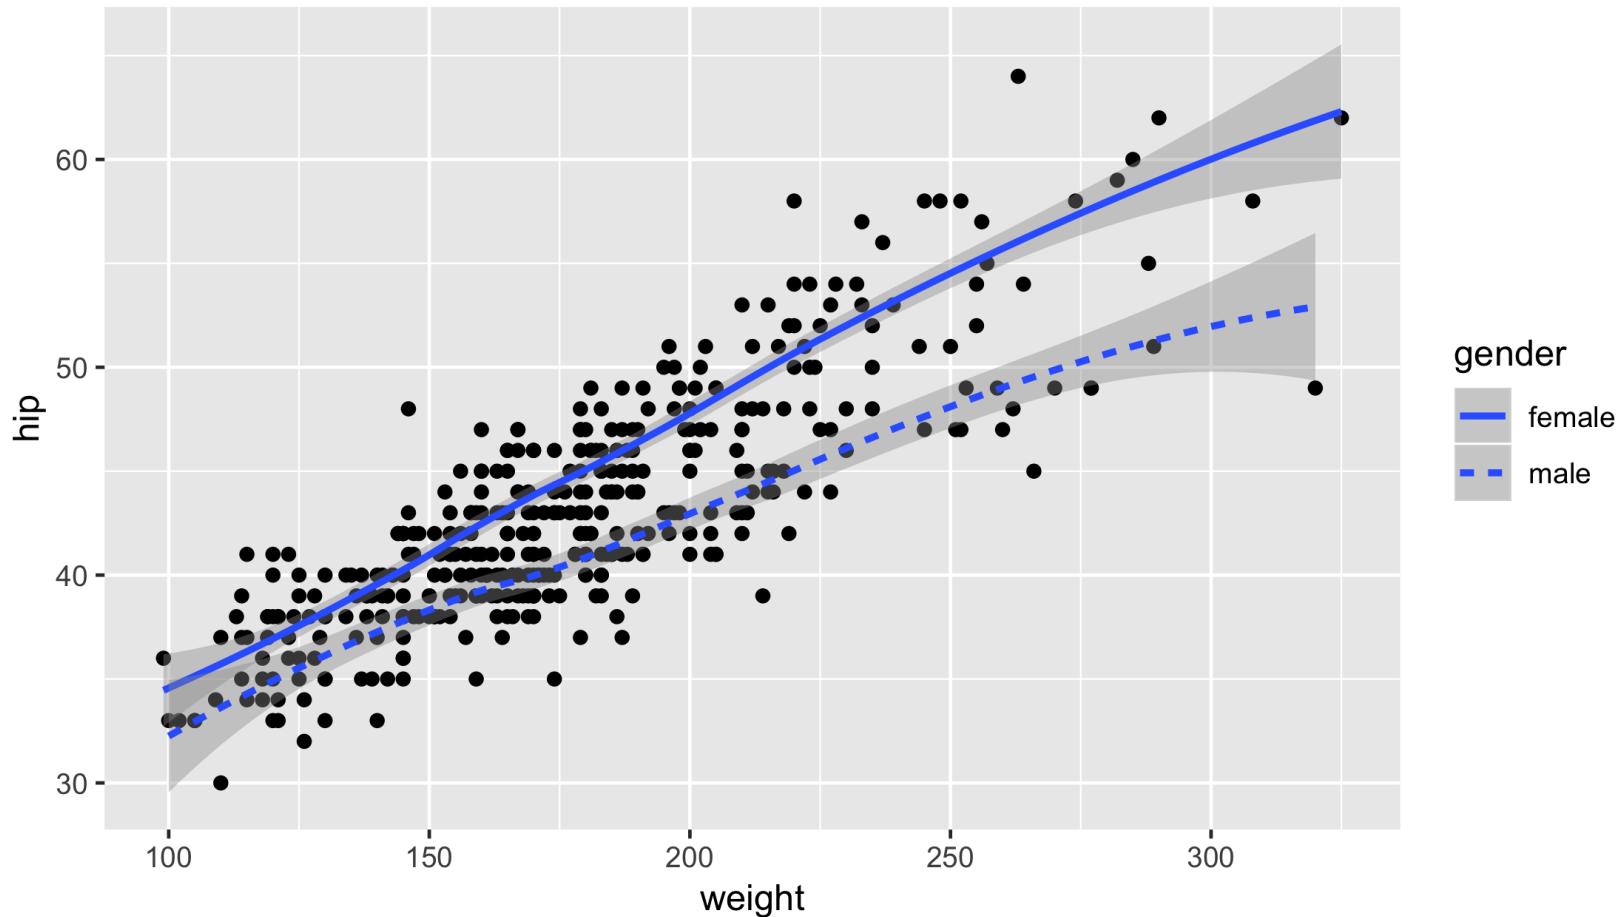
7. Add another layer, theme\_bw(). Remember to use +.

```
ggplot(diabetes, aes(weight, hip)) +  
  geom_point() +  
  geom_smooth()
```

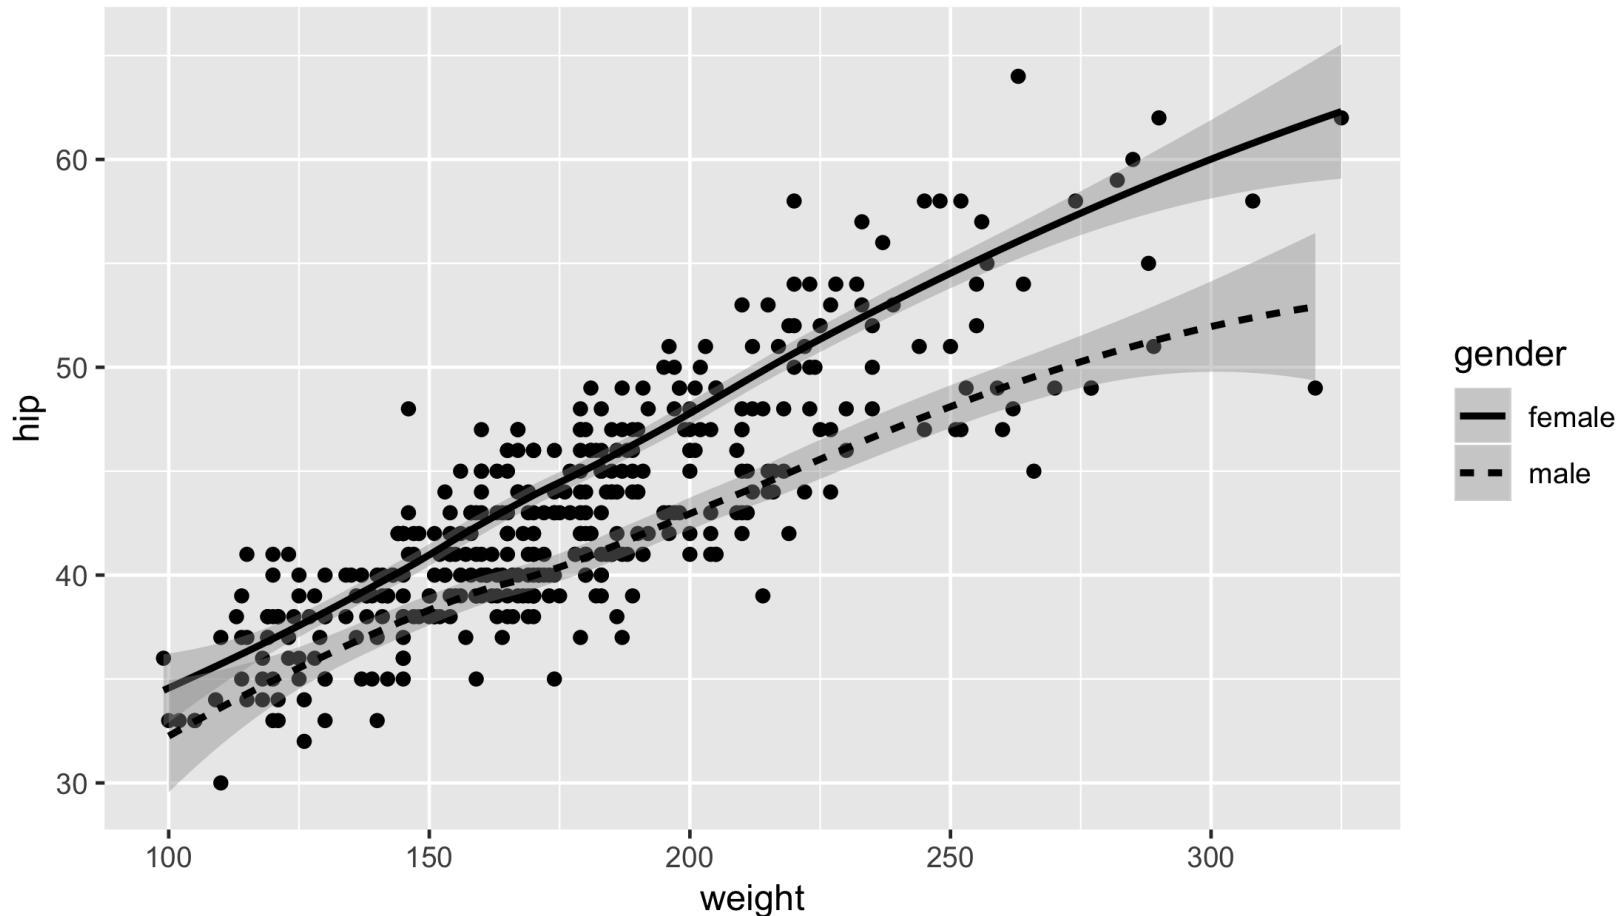
```
ggplot(diabetes, aes(weight, hip)) +  
  geom_point() +  
  geom_smooth()
```



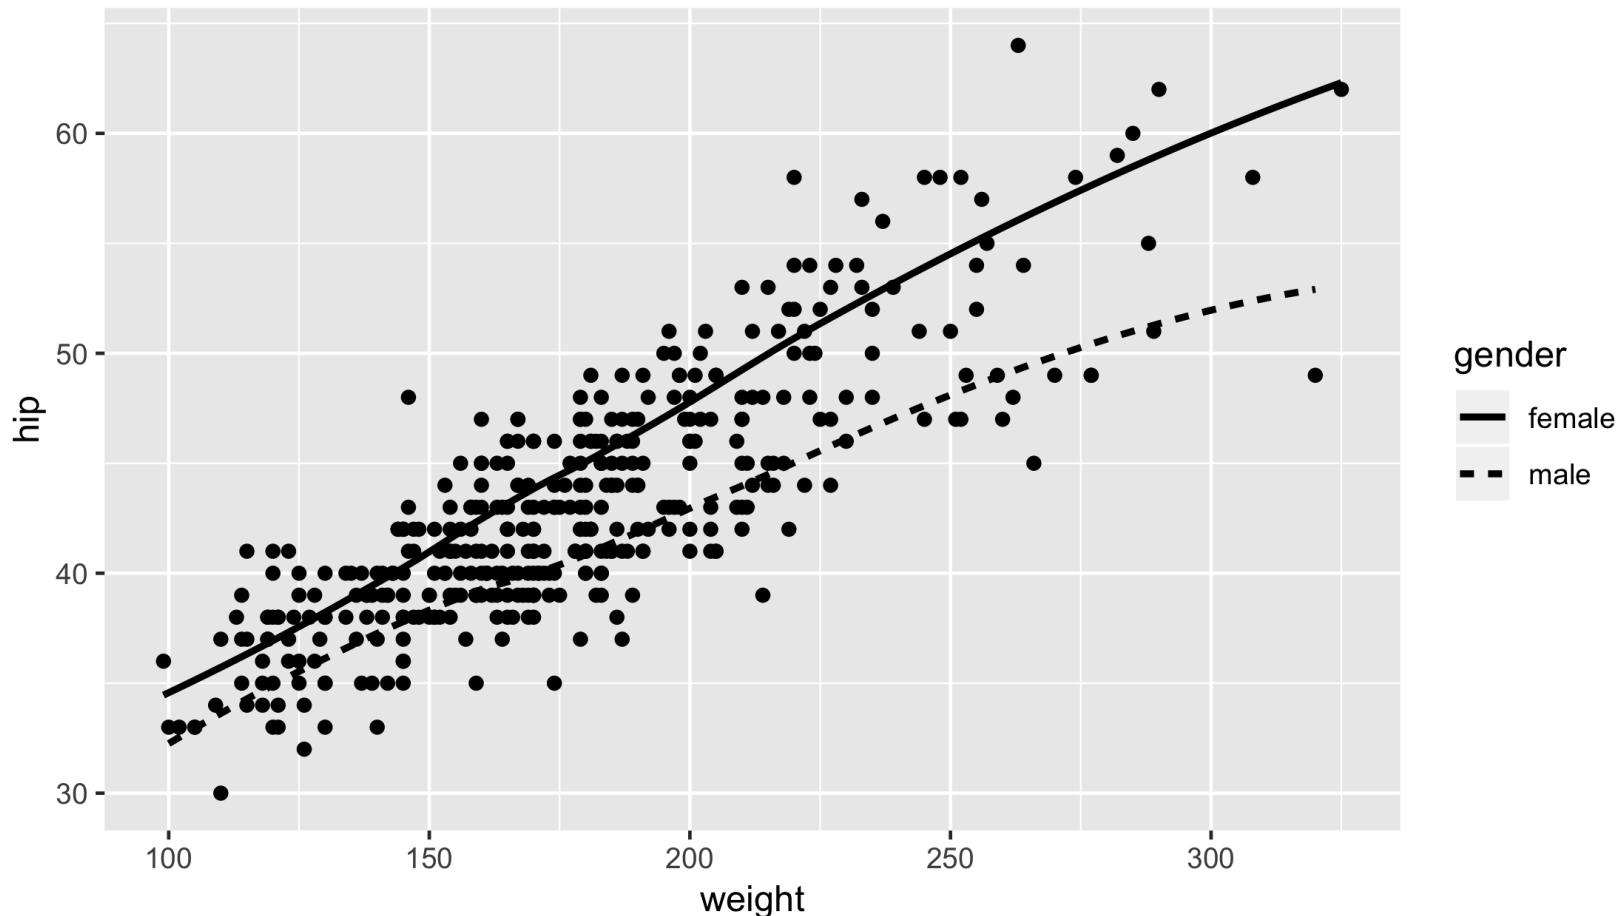
```
ggplot(diabetes, aes(weight, hip)) +  
  geom_point() +  
  geom_smooth(aes(linetype = gender))
```



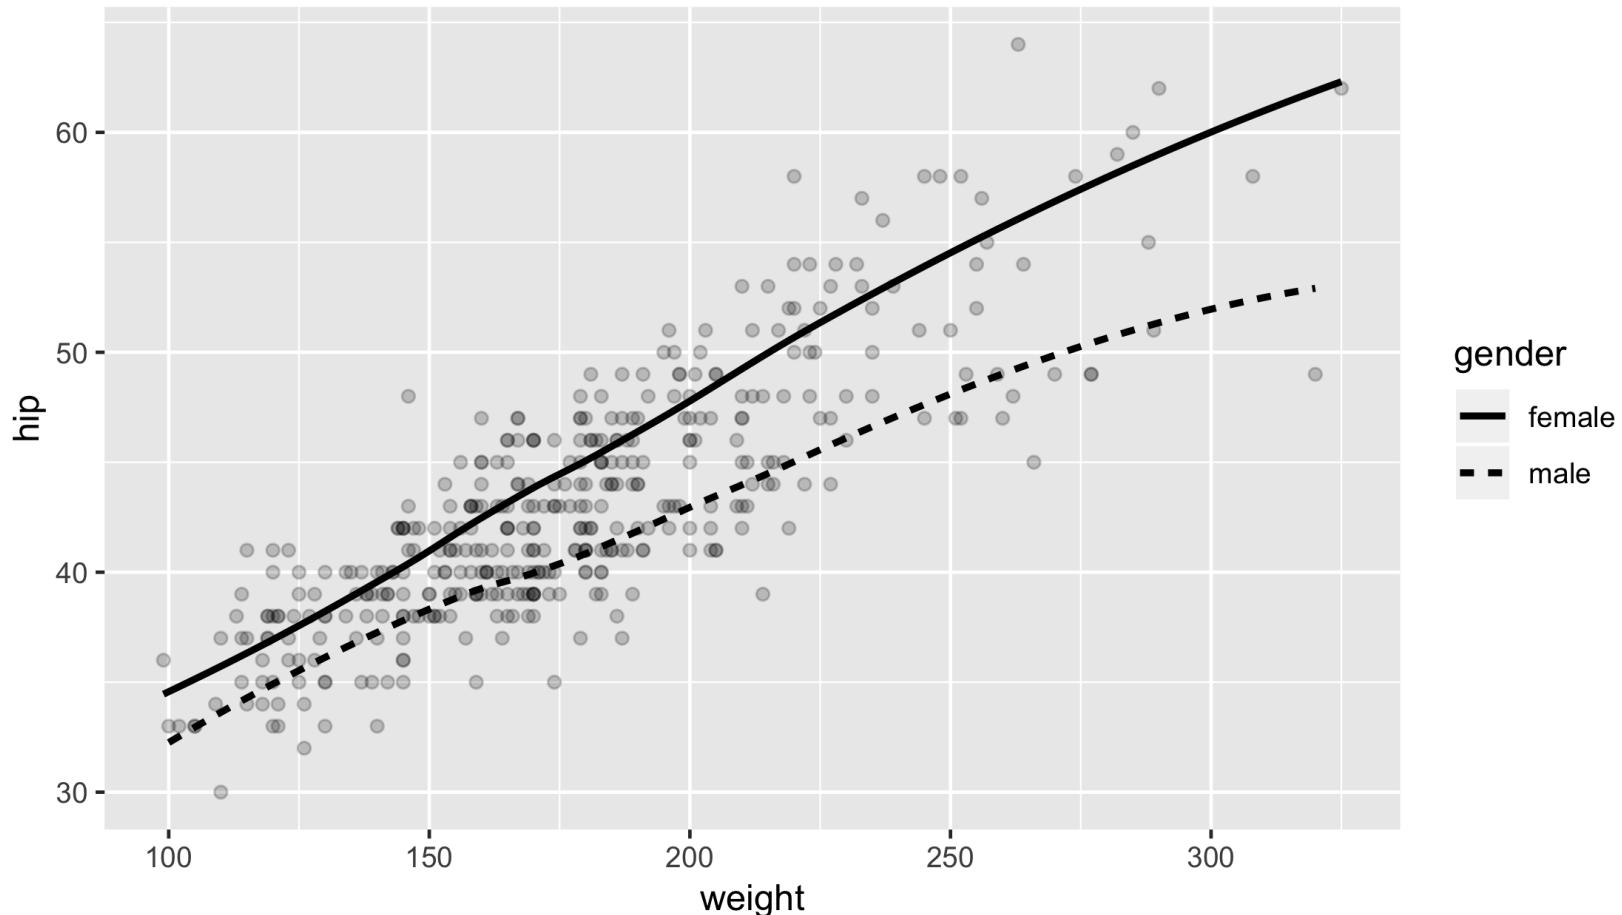
```
ggplot(diabetes, aes(weight, hip)) +  
  geom_point() +  
  geom_smooth(aes(linetype = gender), col = "black")
```



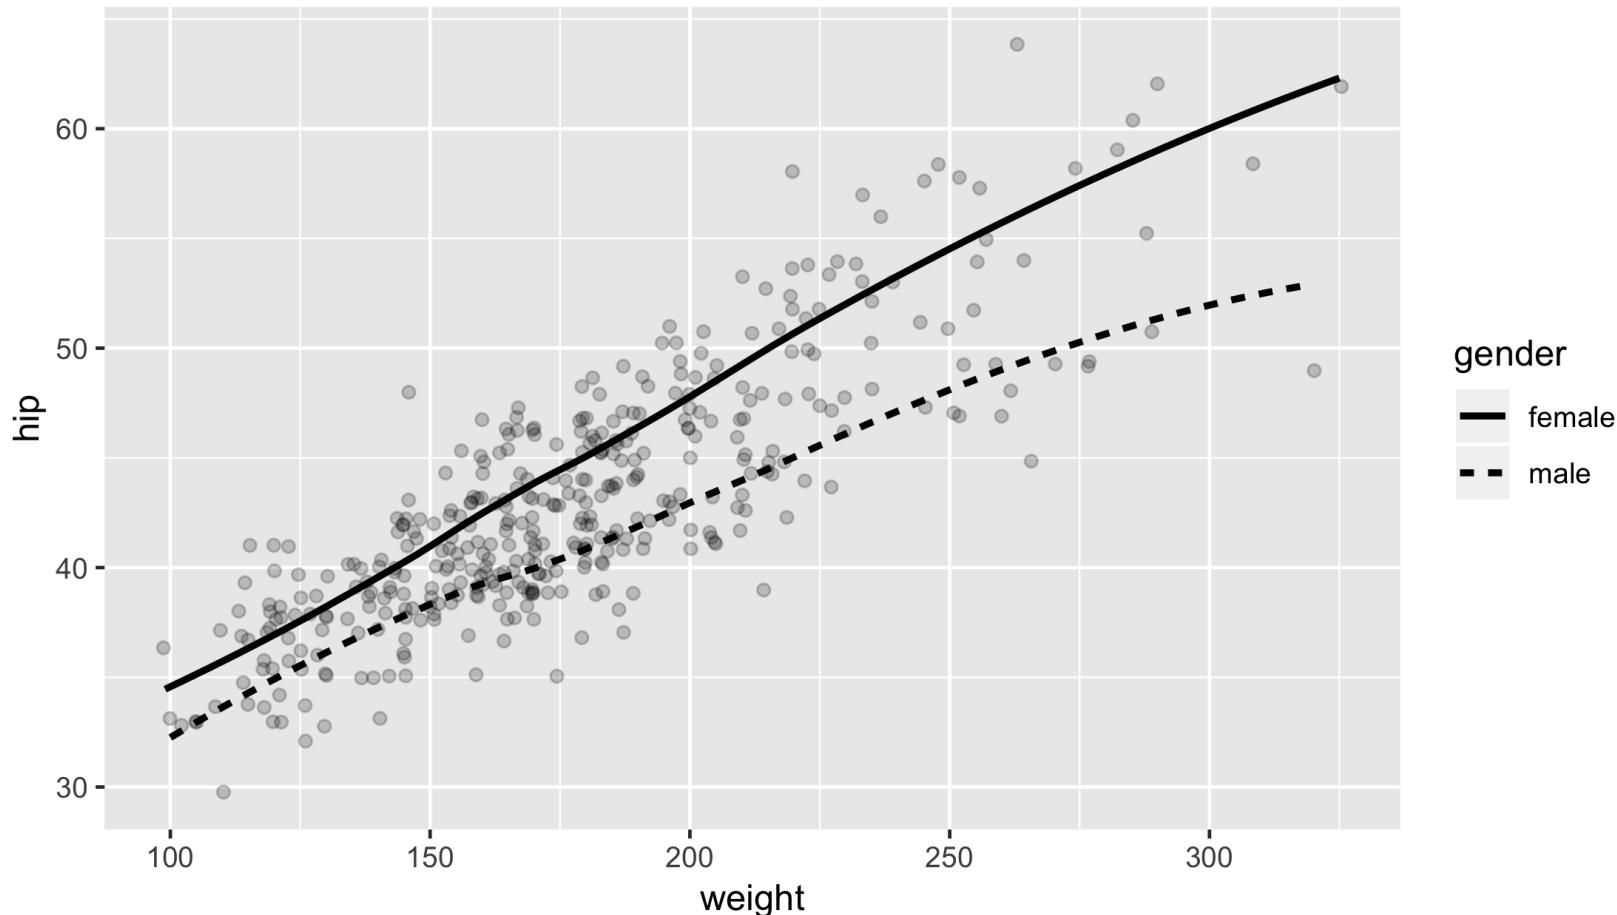
```
ggplot(diabetes, aes(weight, hip)) +  
  geom_point() +  
  geom_smooth(aes(linetype = gender), col = "black", se = FALSE)
```



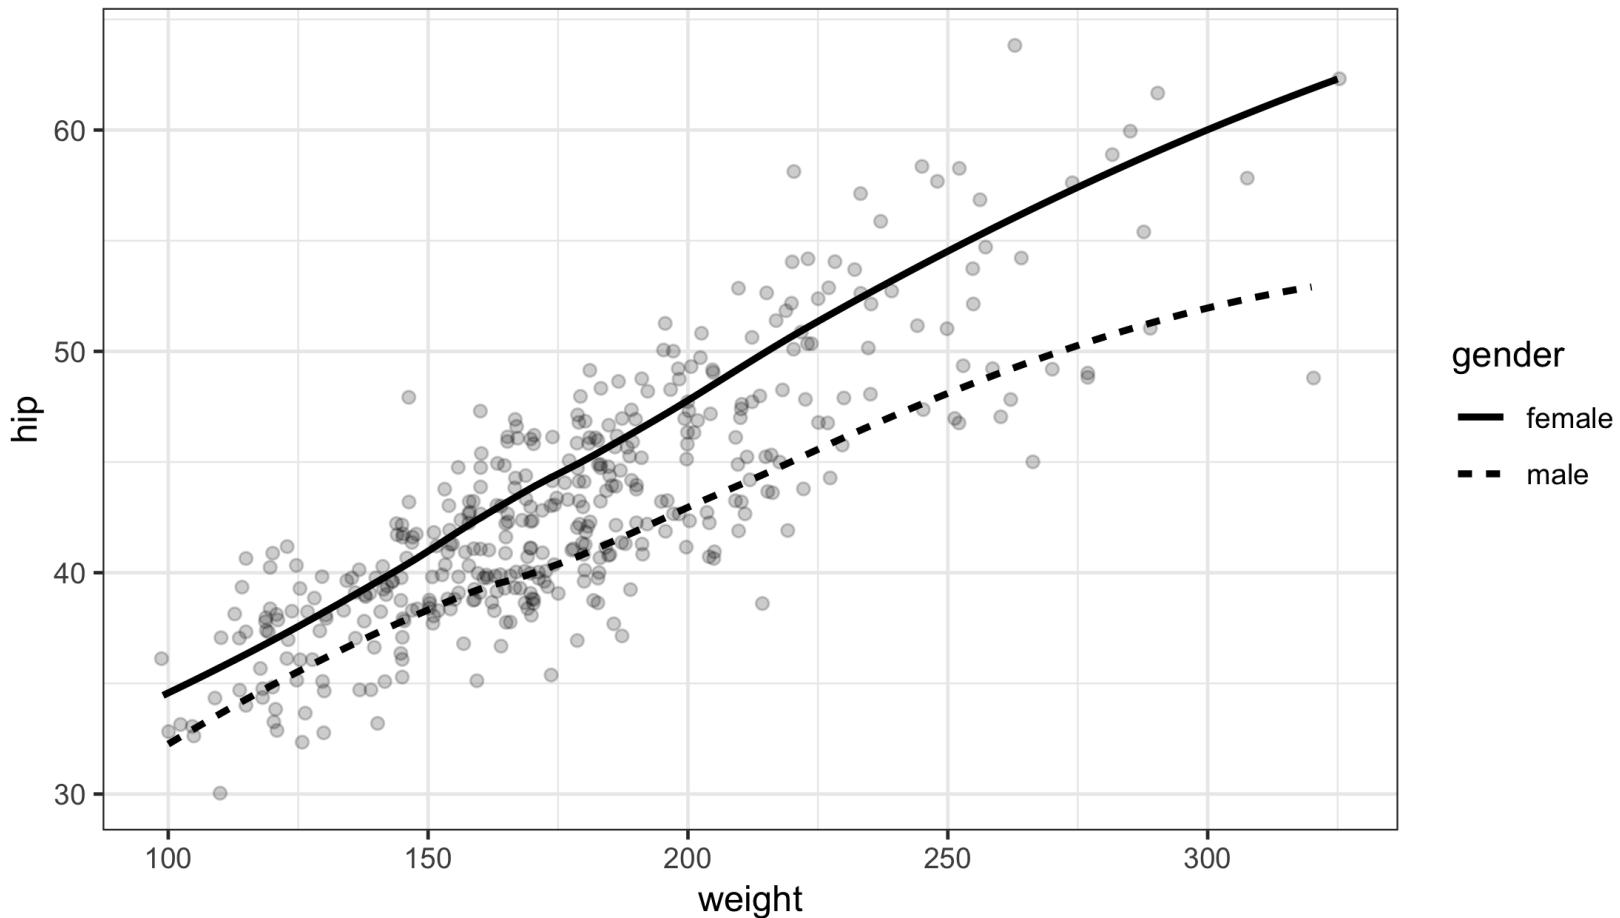
```
ggplot(diabetes, aes(weight, hip)) +  
  geom_point(alpha = .2) +  
  geom_smooth(aes(linetype = gender), col = "black", se = FALSE)
```



```
ggplot(diabetes, aes(weight, hip)) +  
  geom_jitter(alpha = .2) +  
  geom_smooth(aes(linetype = gender), col = "black", se = FALSE)
```



```
ggplot(diabetes, aes(weight, hip)) +  
  geom_jitter(alpha = .2) +  
  geom_smooth(aes(linetype = gender), col = "black", se = FALSE) +  
  theme_bw()
```



# Facets

Easy peazy panels

# Facets

Easy peazy panels

`facet_grid()`

`facet_wrap()`

# Facets

Easy peazy panels

`facet_grid()`

`facet_wrap()`

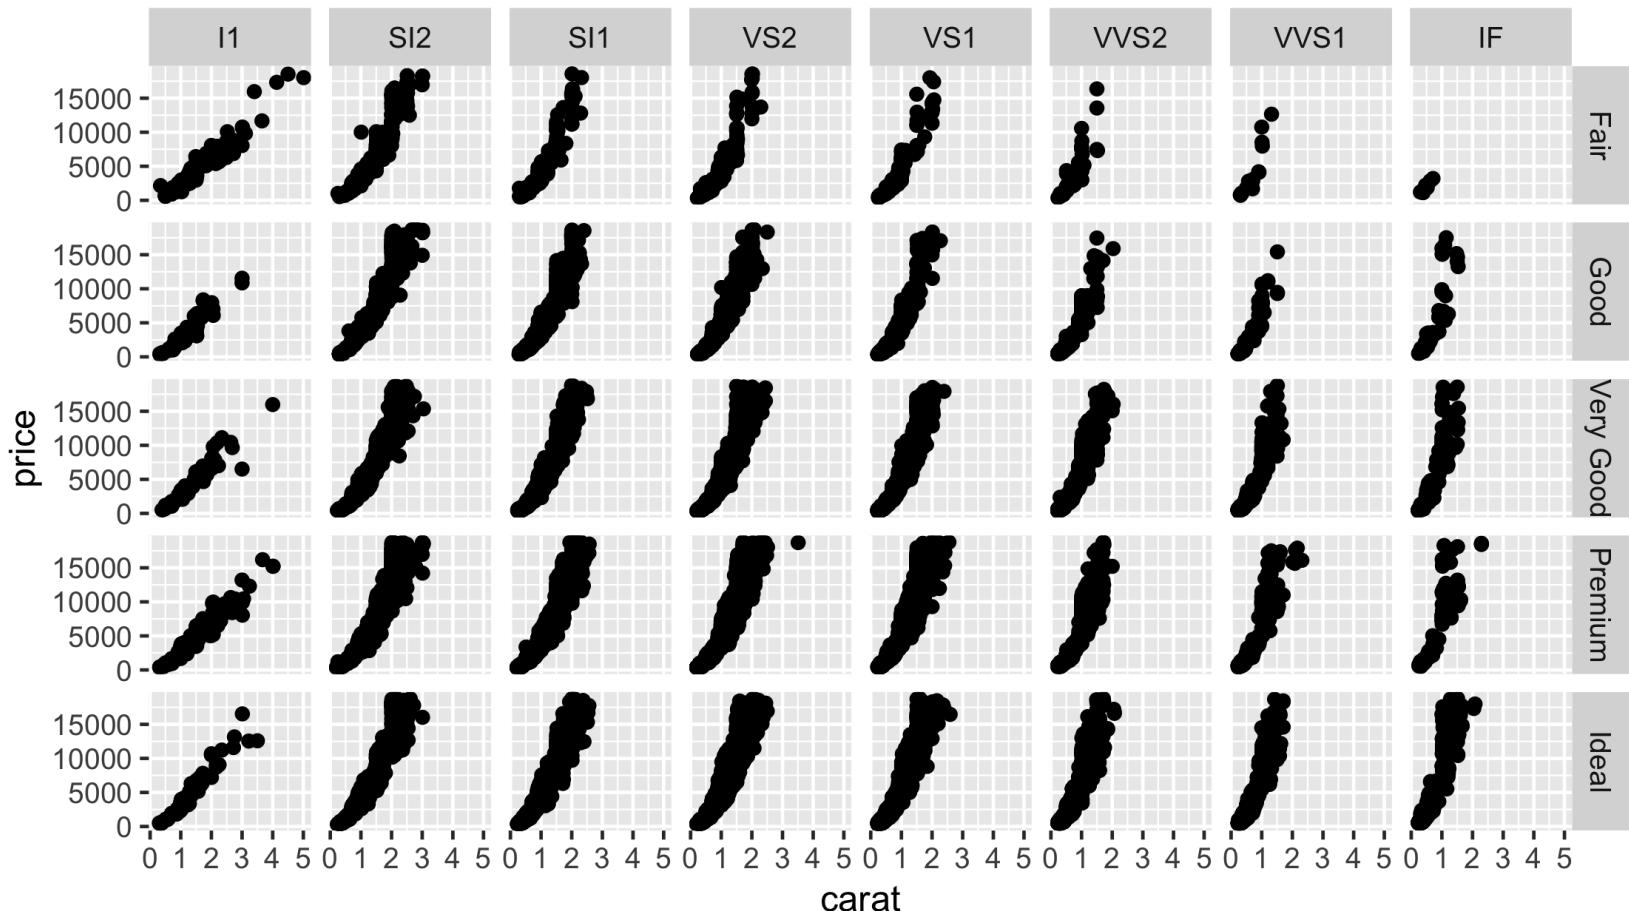
`facet_grid(rows = vars(x), cols = vars(y))`

`facet_wrap(vars(x))`

```
diamonds %>%
  ggplot(aes(x = carat, price)) +
  geom_point() +
  facet_grid(rows = vars(cut), cols = vars(clarity))
```

```
diamonds %>%
```

```
ggplot(aes(x = carat, price)) +  
  geom_point() +  
  facet_grid(rows = vars(cut), cols = vars(clarity))
```

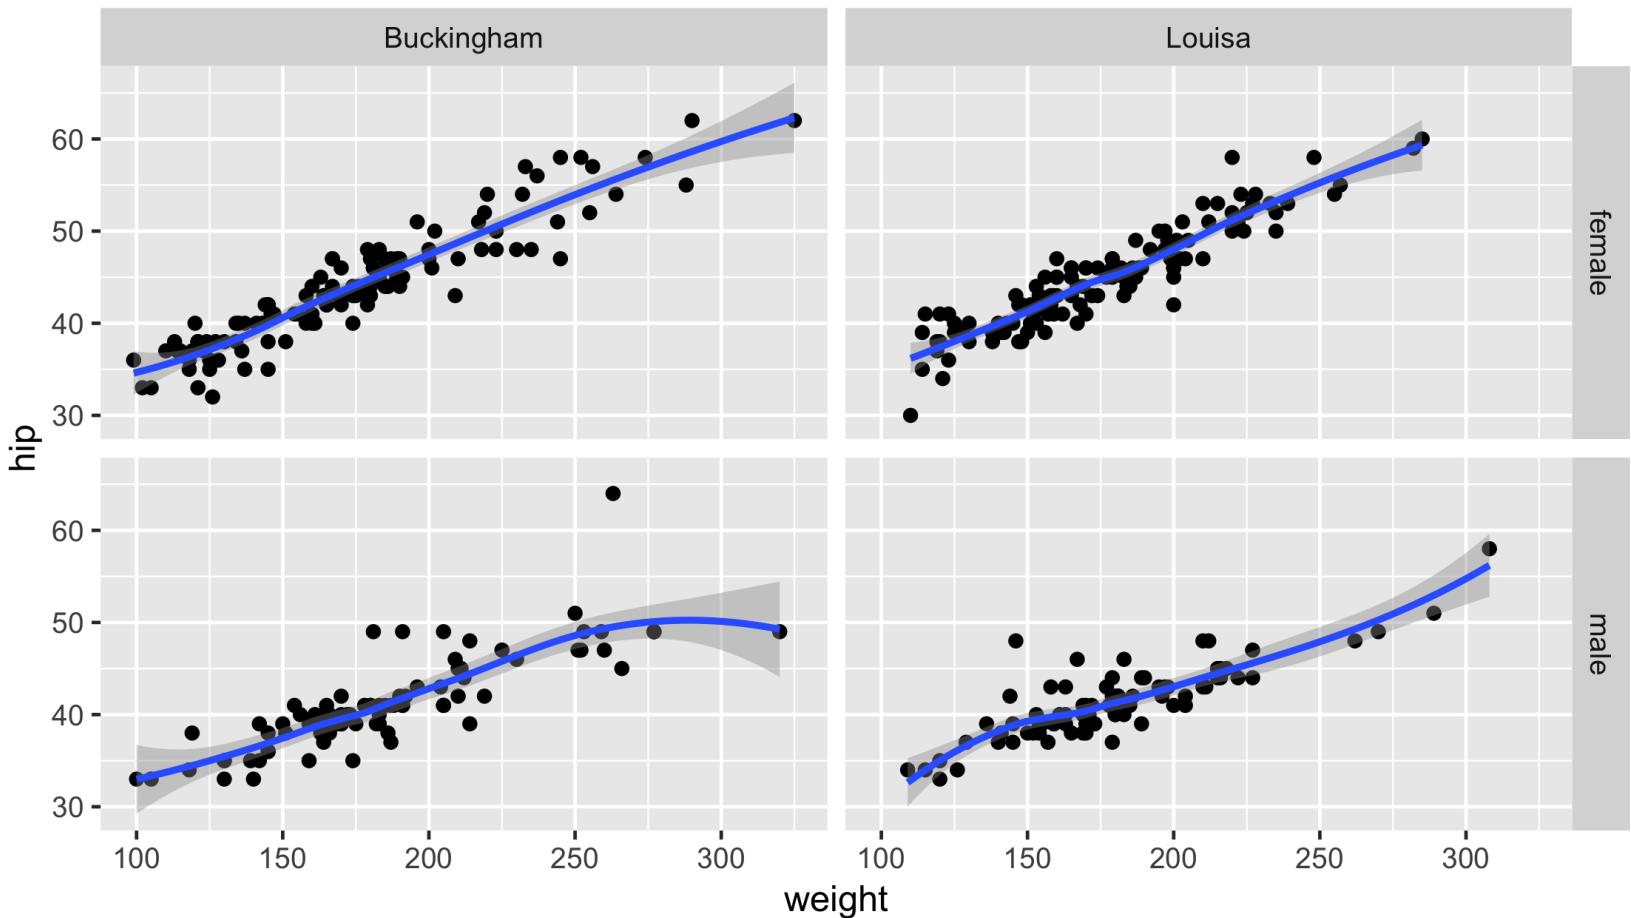


# Your Turn 8

**Use a facet grid by gender and location**

```
ggplot(diabetes, aes(weight, hip)) +  
  geom_point() +  
  geom_smooth()
```

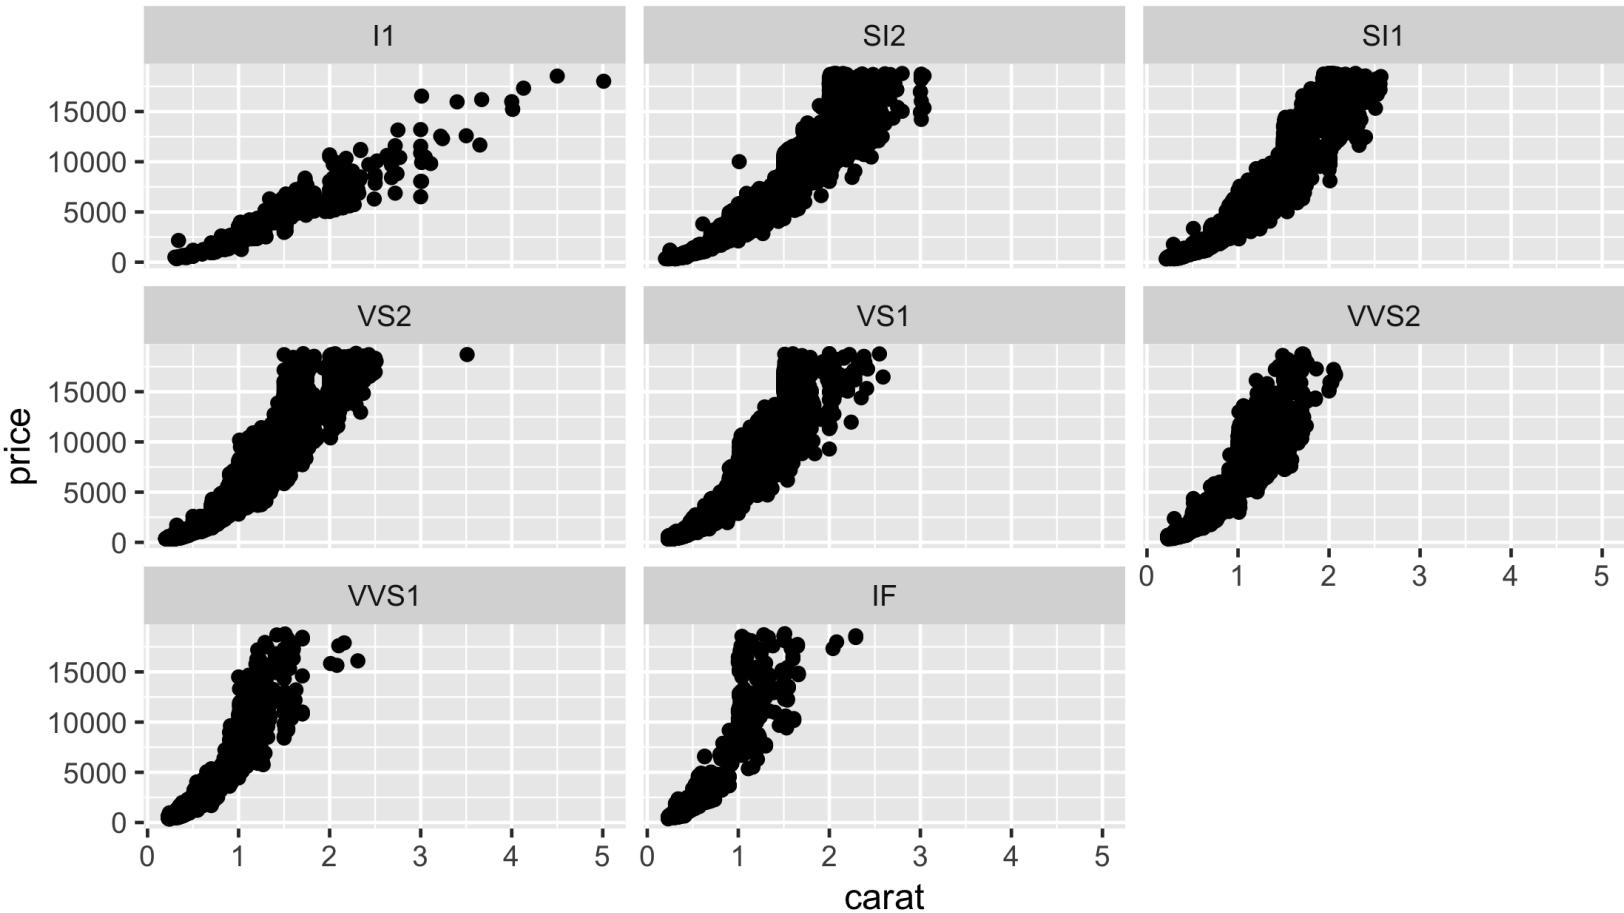
```
ggplot(diabetes, aes(weight, hip)) +  
  geom_point() +  
  geom_smooth() +  
  facet_grid(rows = vars(gender), cols = vars(location))
```



# facet\_wrap()

```
diamonds %>%
  ggplot(aes(x = carat, price)) +
  geom_point() +
  facet_wrap(vars(clarity))
```

# facet\_wrap()



# datasauRus

```
library(datasauRus)
datasaurus_dozen

## # A tibble: 1,846 x 3
##   dataset      x      y
##   <chr>    <dbl> <dbl>
## 1 dino     55.4  97.2
## 2 dino     51.5  96.0
## 3 dino     46.2  94.5
## 4 dino     42.8  91.4
## 5 dino     40.8  88.3
## 6 dino     38.7  84.9
## 7 dino     35.6  79.9
## 8 dino     33.1  77.6
## 9 dino     29.0  74.5
## 10 dino    26.2  71.4
## # ... with 1,836 more rows
```



# new data alert!



## datasaurus\_dozen

	dataset	x	y	z
1	dino	55.38460	97.179500	
2	dino	51.53850	96.025600	
3	dino	46.15380	94.487200	
4	dino	42.82050	91.410300	
5	dino	40.76920	88.333300	
6	dino	38.71790	84.871800	
7	dino	35.64100	79.871800	
8	dino	33.07690	77.564100	
9	dino	28.97440	74.487200	
10	dino	26.15380	71.410300	
11	dino	23.07690	66.410300	
12	dino	22.30770	61.794900	
13	dino	22.30770	57.179500	
14	dino	23.33330	52.948700	
15	dino	25.89740	51.025600	
16	dino	29.48720	51.025600	
17	dino	32.82050	51.025600	
18	dino	35.38460	51.410300	

Where does it come from?

The datasauRus R package

How can I use it?

```
library(datasauRus)  
View(datasaurus_dozen)
```



*it's invisible!*

# Your Turn 9: Challenge!

- 1. Load the `datasauRus` package. This package includes a data set called `datasaurus_dozen`.**
- 2. Use `dplyr` to summarize the correlation between `x` and `y`. First, group it by dataset, and then summarize with the `cor()` function. Call the new variable `corr`. What's it look like?**
- 3. Mutate `corr`. Round it to 2 digits. Then, mutate it again (or wrap it around your first change) using: `paste("corr:", corr)`**
- 4. Save the summary data frame as `corrs`.**
- 5. Pass `datasaurus_dozen` to `ggplot()` and add a point geom**
- 6. Use a facet (wrap) for dataset.**
- 7. Add a text geom. For this geom, set `data = corrs`. You also need to use `aes()` in this call to set `label = corr`, `x = 50`, and `y = 110`.**

```
corrs <- datasaurus_dozen %>%
  group_by(dataset) %>%
  summarize(corr = cor(x, y)) %>%
  mutate(
    corr = round(corr, 2),
    corr = paste("corr:", corr)
  )
```

```
corrs <- datasaurus_dozen %>%
  group_by(dataset) %>%
  summarize(corr = cor(x, y)) %>%
  mutate(
    corr = round(corr, 2),
    corr = paste("corr:", corr)
  )
```

```
corrs <- datasaurus_dozen %>%
  group_by(dataset) %>%
  summarize(corr = cor(x, y)) %>%
  mutate(
    corr = round(corr, 2),
    corr = paste("corr:", corr)
  )
```

```
corrs <- datasaurus_dozen %>%
  group_by(dataset) %>%
  summarize(corr = cor(x, y)) %>%
  mutate(
    corr = round(corr, 2),
    corr = paste("corr:", corr)
  )
```

## corrs

```
## # A tibble: 13 x 2
##   dataset    corr
##   <chr>     <chr>
## 1 away      corr: -0.06
## 2 bullseye  corr: -0.07
## 3 circle    corr: -0.07
## 4 dino      corr: -0.06
## 5 dots      corr: -0.06
## 6 h_lines   corr: -0.06
## 7 high_lines corr: -0.07
## 8 slant_down corr: -0.07
## 9 slant_up  corr: -0.07
## 10 star     corr: -0.06
## 11 v_lines  corr: -0.07
## 12 wide_lines corr: -0.07
## 13 x_shape  corr: -0.07
```

```
datasaurus_dozen %>%
  ggplot(aes(x, y)) +
  geom_point() +
  geom_text(data = corrs, aes(label = corr, x = 50, y = 110)) +
  facet_wrap(vars(dataset))
```

```
datasaurus_dozen %>%
  ggplot(aes(x, y)) +
  geom_point() +
  geom_text(data = corrs, aes(label = corr, x = 50, y = 110)) +
  facet_wrap(vars(dataset))
```

```
datasaurus_dozen %>%
  ggplot(aes(x, y)) +
  geom_point() +
  geom_text(data = corrs, aes(label = corr, x = 50, y = 110)) +
  facet_wrap(vars(dataset))
```

